



Distributed Systems / Middleware Distributed Programming in Erlang

Alessandro Sivieri

Dipartimento di Elettronica e Informazione
Politecnico di Milano, Italy

sivieri@elet.polimi.it

<http://corsi.dei.polimi.it/distsys>

Slides based on previous works by Alessandro Margara



Outline

- **Introduction**
- Sequential Programming
 - Data structures
 - Single assignment
 - Pattern matching
 - Functional abstractions
 - Dynamic code loading
- Concurrent / Distributed Programming
 - Processes
 - Fault tolerance
 - Distributed Erlang
 - Socket based distribution
- OTP Introduction

Why Erlang?

The world is parallel.

If you want to write programs that behave as other objects behave in the real world, then these programs will have a concurrent structure.

Use a language that was designed for writing concurrent applications, and development becomes a lot easier.

Erlang programs model how we think and interact.

Joe Armstrong



Erlang history

- 1982-1986. Programming experiments: how to program a telephone exchange
- 1986. Erlang emerges as a dialect of Prolog. Implementation is a Prolog interpreter
 - 1 developer (Joe Armstrong)
- 1986. Own abstract machine, JAM
 - 3 developers, 10 users
- 1993. Turbo Erlang (BEAM compiler)
- 1993. Distributed Erlang
- 1996. OTP (Open Telecom Platform) formed
- 1996. AXD301 switch announced
 - Over a million lines of Erlang
 - Reliability of nine 9s

Erlang history/2

- 1998. Erlang banned within Ericsson for other products
- 1998. Erlang “fathers” quit Ericsson
- 1998. Open source Erlang
- 2004. Armstrong re-hired by Ericsson
- 2006. Native symmetric multiprocessing is added to runtime system
- 2011. December. Latest stable release: R15B

Getting started

- To run Erlang programs we will use the BEAM emulator
- Similar to Java JVM
 - Programs are compiled in BEAM ByteCode ...
 - ... and then executed inside the emulator
- Similar to Python
 - It offers an interactive shell ...
 - ... that we will use to run our examples
- To start the BEAM compiler type the command **erl**

Installing Erlang

- Windows
 - Binary installation of the latest version are available at <http://www.erlang.org/download.html>
- Linux (Debian-based systems)
 - apt-get install erlang
- Linux / Mac OS X
 - Build from sources
 - Download latest available version (R15B) at <http://www.erlang.org/download.html>
 - Compile and install



What is Erlang?

- Erlang is a *functional* and *concurrent* programming language
- Why functional?
 - Computation is performed by means of mathematical function evaluation
 - Often recursive
 - Functions are first-class values
 - Can be used as parameters to define higher order abstractions
- Why concurrent?
 - Asynchronous message passing
 - Message passing = No shared memory
 - No side effects
 - No locks
 - Asynchronous = No synchronous invocations
 - Isolation between processes
 - Fault-tolerance
 - Efficient concurrency management
 - Lightweight processes and efficient communication

Our approach

- Few slides on the syntax
- Many examples
 - Available online as source code
- Focusing on the following aspects
 - Features and abstractions offered by functional programming languages
 - Concurrent / distributed programming
- We will use only base Erlang
 - We will mention some abstractions built inside existing libraries as examples of functional programming power



Outline

- Introduction
- **Sequential Programming**
 - **Data structures**
 - **Single assignment**
 - **Pattern matching**
 - **Functional abstractions**
 - **Dynamic code loading**
- Concurrent / Distributed Programming
 - Processes
 - Fault tolerance
 - Distributed Erlang
 - Socket based distribution
- OTP Introduction



C quicksort

```
void QuickSort(int list[], int beg, int end)
{
    int piv; int tmp;
    int l,r,p;

    while (beg < end)
    {
        l = beg; p = (beg + end) / 2; r = end;
        piv = list[p];
        while (1)
        {
            while ((l <= r) && ((list[l] - piv) <= 0 )) l++;
            while ((l <= r) && ((list[r] - piv) > 0 )) r--;
            if (l > r) break;
            tmp = list[l]; list[l] = list[r]; list[r] = tmp;
            if (p==r) p=l;
            l++; r--;
        }
        list[p] = list[r]; list[r] = piv;
        r--;
        if ((r - beg) < (end - 1))
        {
            QuickSort(list, beg, r);
            beg = l;
        }
        else
        {
            QuickSort(list, l, end);
            end = r;
        }
    }
}
```



Erlang quicksort

```
qsort([]) -> [];  
qsort([Pivot|Rest]) ->  
    qsort([ X || X <- Rest, X < Pivot]) ++ [Pivot] ++ qsort([ Y || Y <- Rest, Y >= Pivot]).
```

Termination characters syntax...

- ... or: the main error source in your/our first Erlang listings
- Four possible termination characters:
 - ‘.’ is used for single lines in the shell or for the last line of a function
 - ‘,’ is used for each intermediate line in a function
 - ‘;’ is used for terminating a code block inside ***case/if/receive/try/catch*** (more on this later)
 - “ (no termination character) is used for terminating the last code block inside ***case/if/receive/try/catch***
- Thank Prolog for all this mess!

Variables

- Variables must start with a capital letter
- Variables are untyped
 - $A = 123456789$.
 - $B = \text{“erlang”}$.
 - $C = 123456.12 * 654321.345$.
- Variables don't vary!!!
 - Single assignment

Single assignment

- A variable that has had a value assigned to it is called a ***bound*** variable ...
- ... otherwise it is called an ***unbound*** variable
- All variables start off unbound
- When Erlang sees a statement such as $X = 1234$, it binds variable X to the value 1234
- Before getting bound, X could take any value
- Once it gets a value, it holds on to it forever

Single assignment

- Single assignment is like algebra
 - If you use a variable X on different parts of an equation, it always keeps the same value
 - Not like in imperative programming languages where statements like $X = X + 1$ are allowed
 - In Erlang $X = X + 1$ is an error
- Why single assignment is good?
 - Only one possible value inside a given scope
 - Increases readability
 - X will always represent the same value
 - Prevents from modifications in global state
 - Global variables cannot be modified by functions
 - Forces better design choices
 - Isolation of different functions
 - Fault-tolerance
 - Hot-swap



Atoms

- Atoms are used to represent different non-numerical constant values
- Atoms start with a lower case letter
- Atoms are global, and this is achieved without the use of macro definition or include files
- The value of an atom is just the atom
- If you want to write a calendar application the atoms “monday”, “tuesday” etc. will have the same value everywhere
- Sometimes, you may need ‘ (ticks) for specifying atom names with strange characters
- Remember that also module, function, host names (and many other types) are actually atoms

Tuples

- A tuple is a structure composed by a fixed number of unnamed fields
- For example $X = \{\text{temp}, 12\}$
 - creates a tuple with two fields
 - the first field is the atom “temp”
 - the second field is the integer value 12
 - the tuple is assigned to variable X



Pattern matching

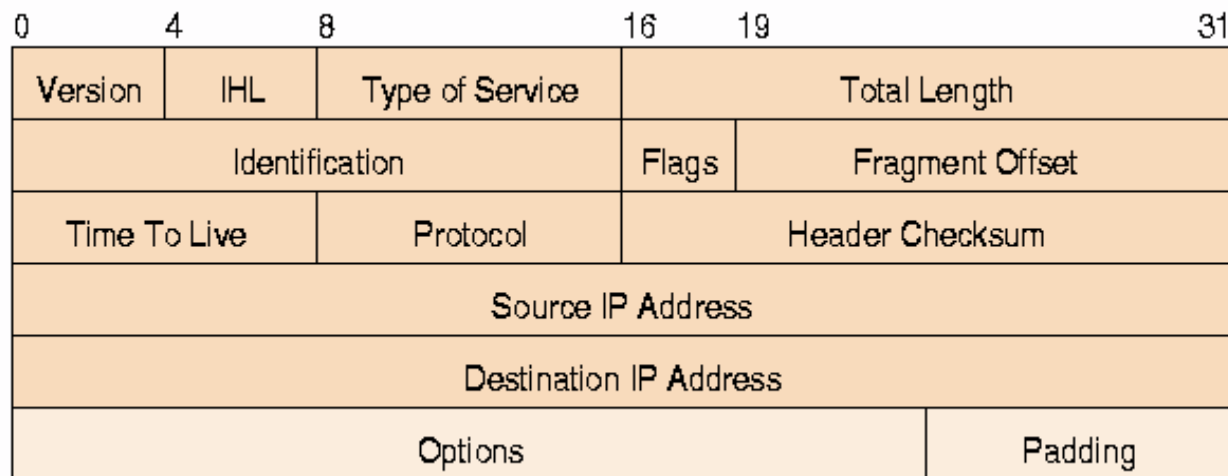
- Pattern matching is a central concept in Erlang
- The pattern matching operator is =
- It evaluates the right side and then matches the result against the pattern on the left side
- We have already seen an example of pattern matching
 - Variable assignment
 - $X = 10$
 - Erlang says to itself “What can I do to make this statement true?”
 - In this case it binds the value 10 to the variable X, so that the equation is satisfied

Pattern matching

- Pattern matching works with tuples ...
- ... enabling programmers to extract values
- $\text{Point} = \{\text{point}, 10, 12\}$
- $\{\text{point}, X, Y\} = \text{Point}$
 - Assigns 10 to variable X and 12 to variable Y
- $\{\text{point}, Z, Z\} = \text{Point}$
 - Returns an error: it is not possible to make the statement true
- $\{_, _, W\} = \text{Point}$
 - Assigns 12 to variable W
 - $_$ is the anonymous variable: different occurrences don't have to hold the same value

Bit syntax

- Erlang has a very interesting bit syntax
 - Header = <<IpVersion:4, HLen:4, Srvctype:8, TotLen:16, ID:16, Flgs:3, FragOff:13, TTL:8, Proto:8, HdrChkSum:16, SrcIP:32, DestIP:32, RestDgram/binary>>



Pattern matching on bits

- You can specify:
 - The type of the variable (integer, float, binary)
 - The signedness
 - The endianness
 - The size (from 1 to 256 bits)
- You can match some pattern in *case/receive...*
- Remember to add a final field ready to get the rest of the variable, especially if working on network protocols
 - the payload has (probably) no fixed size



Lists

- Lists are used to store multiple things
 - Es. ToBuy = [{apple, 10}, {pear, 12}, {lemon, 3}]
- Lists can have heterogeneous elements
 - Es. [erlang, 10, {lemon, 3}]
- The first element of a list is called ***Head***
- If you remove the Head from the list what's left is called the ***Tail*** of the list
- If T is a list than also [H | T] is a list
 - | separates the Head from the Tail
- [] is the empty list



Lists

- Pattern matching can be applied to lists as well
- $\text{Buy} = [\{\text{apple}, 10\}, \{\text{pear}, 12\}, \{\text{orange}, 4\}, \{\text{lemon}, 6\}]$
- $[\text{AppleTuple}, \text{PearTuple} \mid \text{Others}] = \text{Buy}$
 - Assigns $\{\text{apple}, 10\}$ to AppleTuple
 - Assigns $\{\text{pear}, 12\}$ to PearTuple
 - Assigns $[\{\text{orange}, 4\}, \{\text{lemon}, 6\}]$ to Others
- $\text{NewBuy} = [\{\text{milk}, 2\} \mid \text{Others}]$
 - Assigns $[\{\text{milk}, 2\}, \{\text{orange}, 4\}, \{\text{lemon}, 6\}]$ to NewBuy

List comprehensions

- Creating lists from existing lists
- Standard construct in functional languages
- [Element | Element <- List, *conditions*]
- Example:

```
> NewList = [X || X <- [1,2,a,3,4,b,5,6], integer(X), X > 3].  
[4,5,6]
```



Strings

- Strictly speaking there are no strings in Erlang
- Strings are really just lists of integer
- Strings are enclosed in double quotation marks
 - For example you can write `S = "Hello"`
 - `"Hello"` is just a shorthand for the list of integers representing the individual characters in the string
- The shell prints a list of integers as a string if and only if all integers in the list represent a printable character

Dictionaries

- A dictionary performs exactly as a list of 2-dimension tuples
- `[{Key1, Value1}, {Key2, Value2}, {Key3, Value3}]`
- Becomes:
 - `dict:new()`
 - `dict:append(NewKey, NewValue, Dict)`
 - `Value = dict:fetch(Key, Dict)`
 - ...
- You can move from lists to dictionaries and back easily (`dict:from_list`, `dict:to_list`)



Modules

- Erlang programs are splitted in modules
- Each module is a “.erl” file
- To compile a module m.erl you can either
 - Call erlc m.erl from outside the BEAM emulator
 - Or call c(m) from within the BEAM emulator
- This creates a m.beam file, containing the bytecode of the module

Modules

- Each module consists of a set of functions
 - Used internally
 - Or externally visible
- Each module starts with
 - `-module(module-name).`
 - `module-name` must be the name of the file
 - `-export([fun1/arity1, fun2/arity2, ... fun-n/arity-n])`
 - Where `fun1`, `fun2`, ... `fun-n` are the names of the functions that have to become visible outside the module
 - And `arity1`, `arity2`, ... `arity-n` are the arity (i.e. number of input parameters) required by each function



Functions

- A function is univocally identified by a name and an arity
- Each function consists of an ordered list of clauses
- Each clause has a pattern and a piece of code
- During a function call clauses are evaluated in order
- When the pattern of a clause is matched, then the associated code is evaluated and the function returns
- If no single pattern can be matched, then an error is generated



Functions

- Example: we define a module geometry with only one (exported) function of arity 1, which computes the area of different figures
- We can use it within the BEAM emulator to compute the area of a square
 - `geometry:area({square, 10})`

```
-module(geometry).  
-export([area/1]).
```

```
area({square, X}) ->  
    X*X;
```

```
area({rectangle, X, Y}) ->  
    X*Y;
```

```
area({circle, R}) ->  
    3.14*R.
```



Functions

- Sometimes it is useful to check constraints on input values
- For this reason Erlang introduces guards
- Introduced after a pattern, using the ***when*** keyword
- The example shows a ***single guard*** ($X > Y$)
- It is possible to combine single guards using logical ***and*** (,) or logical ***or*** (;)
- Beware that if you want short-circuit expressions, you have to use ***andalso*** or ***orelse***

```
-module(guards).  
-export(max/2).
```

```
max(X, Y) when X > Y ->  
    X;  
max(X, Y) -> Y.
```




Case and if expressions

- Pattern matching and guards can be used to define conditional blocks
 - *Case* expressions
 - *If* expressions
- Expressions are evaluated in order
- If no match is found an error is generated
- Beware that in *If* expressions you always need the “else” part

```
case Expression of
  Pattern1 [when Guard1] -> Expr_seq1;
  Pattern2 [when Guard2] -> Expr_seq2;
  ...
  Any -> io:format("Unknown sequence: ~p~n", [Any])
end
```

```
if
  Guard1 ->
    Expr_seq1;
  Guard2 ->
    Expr_seq2;
  ...
  true ->
    Default_seq
end
```

Erlang: a functional PL

- In just a few slides we have already seen all the building blocks of Erlang
 - We can now write every sequential program
 - Without *while*, *for* statements!
- Erlang is a functional programming language
 - Everything is performed through function evaluation
 - Functions are values
 - It is possible to assign functions to variables ...
 - ... to use functions as parameters for other functions ...
 - ... and to return function as result of other functions

Erlang: a functional PL

- There are no loop statements
- Iteration is performed using recursive functions
- Example: we want to sum all the elements of a list of integers
- We recursively sum the head to the rest of the list until we arrive to the empty list

```
-module(listSum).  
-export([sum/1]).
```

```
sum([]) ->  
    0;  
sum([H|T]) ->  
    H+sum(T).
```

Saving space: tail recursion

- In the previous code $H + \text{sum}(T)$ cannot be evaluated until the function $\text{sum}(T)$ returns
 - Every function call requires stack space
 - The function $\text{sum}(X)$ evaluates in $O(\text{length}(X))$ space
- We can implement the same function to evaluate in constant space
 - Using an accumulator
 - Using tail recursion (= the last thing a function does is calling itself)
 - Same cost as in imperative programming loops
- Every recursive function can be transformed in a tail recursive function
 - It is good practice to use tail recursion

```
-module(tail).  
-export([tailSum/1]).
```

```
tailSum(X) ->  
    tailSum(X, 0).
```

```
tailSum([H|T], Acc) ->  
    tailSum(T, Acc+H);  
tailSum([], Acc) ->  
    Acc.
```

Higher order function

- A common task is the execution of the same transformation on all the elements of a list
- We can write a single function for each possible transformation
- Or we can use the possibility to use functions as values
 - Map executes a “generic” task on all the elements of a list
 - It is said to be a higher order function

```
-module(map).  
-export([map/2,double/1]).
```

```
double(N) -> N*2.
```

```
map([H|T], F) -> [F(H)|map(T, F)];  
map([], _) -> [].
```

```
>c(map).  
>D = fun(X) -> map:double(X) end.  
>A = [1,2,3].
```

Compile

```
>map:map(A, D).  
>[2,4,6]
```

Assign a function
to a variable

Use the function
as parameter

Functions that return functions

- Not only can functions be used as arguments to functions ...
- ... but functions can also **return** functions
 - It is not used that often, at least wrt to the previous mode
- Suppose we have a list of something (es. Fruit)
- `Fruit = [apple, pear, orange]`
- We can define a function `Test` that returns a function that checks whether an element is in a list
 - `Test = fun(L) -> (fun(X) -> lists:member(X, L) end) end.`
 - `lists:member` is a function that returns true if `X` is in `L`
- We can now create a function `IsFruit`
 - `IsFruit = Test(Fruit)`
 - `IsFruit(apple)` will return true
 - `IsFruit(cat)` will return false

Programming abstractions

- Using higher order functions enables programmers to create different levels of abstractions
- This is conceptually similar to the creation of object hierarchies in Object Oriented Languages like Java or C#
- Object Oriented Languages simplify reuse of code by defining abstract members
- Functional Languages use function parameterization:
 - Functions are values
 - Functions can be used as parameters

Data manipulation

- Functional languages allow you to write extremely compact code when manipulating data
 - `map()`: applies a given function to all members of a list
 - `fold()`: applies a given function to all members of a list, passing an accumulator for getting a result of that function
 - `filter()`: filters a list according to a given function
 - `zip()`: puts together two lists in a single list of 2-dimension tuples
 - ...
- The same holds for dictionaries

Data manipulation

```
List = [1, 2, 3, 4, 5, 6, 7],  
lists:map(fun(Element) -> 2 * Element end, List),  
Sum = lists:foldl(fun(Element, AccIn) -> Element + AccIn, 0, List),  
EvenList = lists:filter(fun(Element) when Element rem 2 == 0 -> true;  
                        (Element) -> false end, List),  
Days = [monday, tuesday, wednesday, thursday, friday, saturday, sunday],  
DayWithNumber = lists:zip(Days, List).
```

Exception handling

- In Erlang exceptions are raised automatically when the system encounters an error
 - Pattern matching errors
 - Function call with incorrectly typed arguments
- It is also possible to throw exceptions explicitly
 - `throw(Why)` throws an exception that the caller is expected to handle
 - `erlang:error(Why)` is used to denote “crashing errors”; something that the caller is not supposed to manage
 - `exit(Why)` explicitly stops a process; if the exception is not managed a message is broadcast to all linked processes (more on this later)

Exception handling

- Exception handling is very similar to a case expression
- ExceptionType is an atom, which defines the kind of exception (throw, exit, error) one wants to catch

```
try FuncOrExpressionSequence of
  Pattern1 [when Guard1] -> Expressions1;
  Pattern2 [when Guard2] -> Expressions2;
  ...
catch
  ExceptionType: ExPattern1 [when ExGuard1] -> ExExpressions1;
  ExceptionType: ExPattern2 [when ExGuard2] -> ExExpressions2;
  ...
after
  AfterExpressions
end
```

Result evaluated here
if no exception occurred

Exceptions (if any)
Evaluated here



Some examples

- Write a function that returns the maximum element in a list

```
-module(max).  
-export([max/1]).
```

```
max([Head|Tail]) ->  
    max(Tail, Head).
```

```
max([], Max) ->  
    Max;
```

```
max([Head|Tail], Max) when Head > Max ->  
    max(Tail, Head);
```

```
max([_|Tail], Max) ->  
    max(Tail, Max).
```

Works only with
non-empty lists!



Some examples

- Write a function that reverses the order of a list

```
-module(reverse).  
-export([reverse/1]).
```

```
reverse(List) ->  
    reverse(List, []).
```


```
reverse([Head | Rest], ReversedList) ->  
    reverse(Rest, [Head|ReversedList]);  
reverse([], ReversedList) ->  
    ReversedList.
```

Dynamic code loading

- Dynamic code loading is a feature directly built inside Erlang
- A module can access a function of another modules in two ways:
 - Importing modules
 - `-import(module_name)`
 - Using fully qualified names
 - `module_name:function_name`
- The former continues to adopt the previously loaded version of a module
- The latter ensures that the latest version of the module is used
 - Even if the module has been recompiled

Dynamic code loading

- Two possible versions of a module can exist at the same time
- No more than two versions are allowed
- If a third version is created:
 - B1 is removed
 - Existing computation aborted
 - B2 continues to exist
 - B3 is the new current version



Dynamic code loading

- Dynamic code loading is a low level feature
- It enables programmers to change system code at runtime
 - To fix bugs
 - To include new functionalities
 - To improve performance
- Higher level abstractions have been designed on top of it
 - OTP (Open Telecom Platform) offers
 - Implementation of design patterns that simplify error-free code loading
 - Tools to automatize installation of new software versione involving multiple modules upgrades



Dynamic code loading: example

```
-module(dynCode1).  
-export([start/0]).
```

```
start() ->  
    spawn(fun loop/0).
```

```
loop() ->  
    Val = dynCode2:val(),  
    io:fwrite("Val = ~p~n", [Val]),  
    sleep(2000),  
    loop().
```

```
sleep(Time) ->  
    receive  
        after Time -> ok  
    end.
```

```
-module(dynCode2).  
-export([val/0]).
```

```
val() ->  
    1.
```

If we change value
and compile, dynCode1
will print the new value

Prints the value
computed by
dynCode2
every 2 seconds



Outline

- Introduction
- Sequential Programming
 - Data structures
 - Single assignment
 - Pattern matching
 - Functional abstractions
 - Dynamic code loading
- **Concurrent / Distributed Programming**
 - **Processes**
 - **Fault tolerance**
 - **Distributed Erlang**
 - **Socket based distribution**
- OTP Introduction

Concurrent programming

- Erlang is sometimes described as a ***concurrency oriented*** programming language ...
- This does not only mean that writing concurrent programs is
 - Possible
 - Easy
 - Efficient
- This refers to the possibility to take into account concurrency when designing complex systems
- Erlang forces programmers to think about processes as independent actors ...
- ... communicating only through message passing

Concurrent programming

- Concurrent programming requires just three new primitives:
 - ***Pid = spawn(Fun)***. Creates a new concurrent process that evaluates *Fun*. The new process runs in parallel with the caller. *Spawn* returns a ***Pid*** (process identifier) which can be used to send messages to the process
 - ***Pid ! Message***. Sends ***Message*** to the process with identifier ***Pid***. Message sending is asynchronous: the sender has not to wait, but can continue its own task
 - ***Receive ... end***. Used to receive messages: messages are evaluated using pattern matching. Messages are stored in a sort of mailbox (persistent!) until the received function is called



Simple Server

```
-module(server).  
-export([start/0]).
```

```
loop() -> receive  
    {Sender, Fun, Num} ->  
        Sender ! Fun(Num),  
        loop()  
end.
```

```
start() -> spawn(fun loop/0).
```

```
> Dup = fun(X) -> 2*X end.  
> c(server).  
> Pid = server:start().  
> Pid ! {self(), Dup, 256}.  
> receive A -> A end.  
512
```

The server waits for messages
containing the sender Pid,
a function and
a number

Sends back Fun(Num)
and waits again

In the shell we start the server

We send a request

We receive the response



Clock

- In this example we define a process that executes a function *Fun* periodically (period = *Time*)
- We introduce two new keywords
 - ***After***: defines what to do if no matching message is received after *Time* elapsed
 - ***Register***: associate the Pid of a process to an atom

```
-module(clock).  
-export([start/2, stop/0]).  
  
start(Time, Fun) ->  
    register(clock, spawn(fun() -> tick(Time, Fun) end)).  
  
stop() -> clock ! stop.  
  
tick(Time, Fun) ->  
    receive  
        stop ->  
            void  
    after Time ->  
        Fun(),  
        tick(Time, Fun)  
    end.
```



Variable

- Write a program that simulates a simple integer variable (allowing init, add, sub and get operations)

```
-module(var).  
-export([init/0, add/1, sub/1, get/0]).
```

```
init() ->  
    Pid = spawn(fun() -> loop(0) end),  
    register(var, Pid).
```

```
loop(N) ->  
    receive  
        {add, X} -> loop(N+X);  
        {sub, X} -> loop(N-X);  
        {Pid, get} -> Pid ! N,  
                    loop(N)  
    end.
```

```
add(X) -> var ! {add, X}.
```


```
sub(X) -> var ! {sub, X}.
```

```
get() -> var ! {self(), get},  
        receive Result -> Result end.
```



Fault tolerance

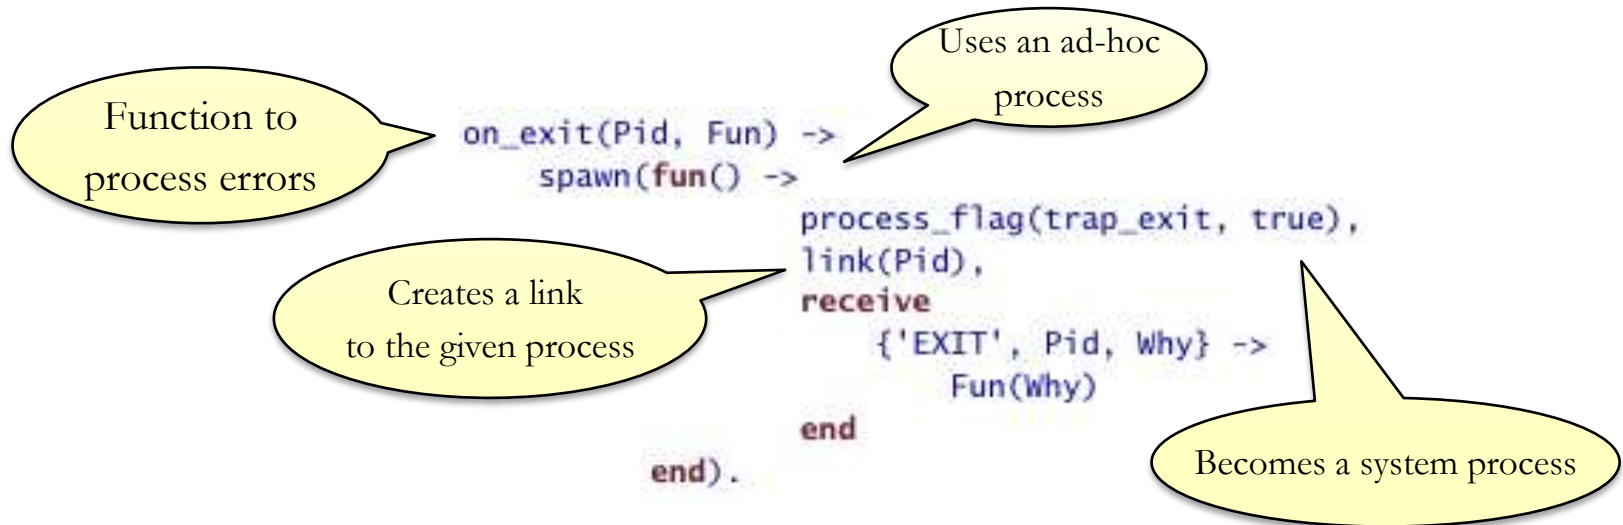
- A key feature of Erlang is its ability to simplify the design of fault tolerant programs
- This is achieved through process linking
 - A process P can link to process Q by calling the *link*(Q) function
 - When one process dies, an *Exit* signal is sent to every linked process



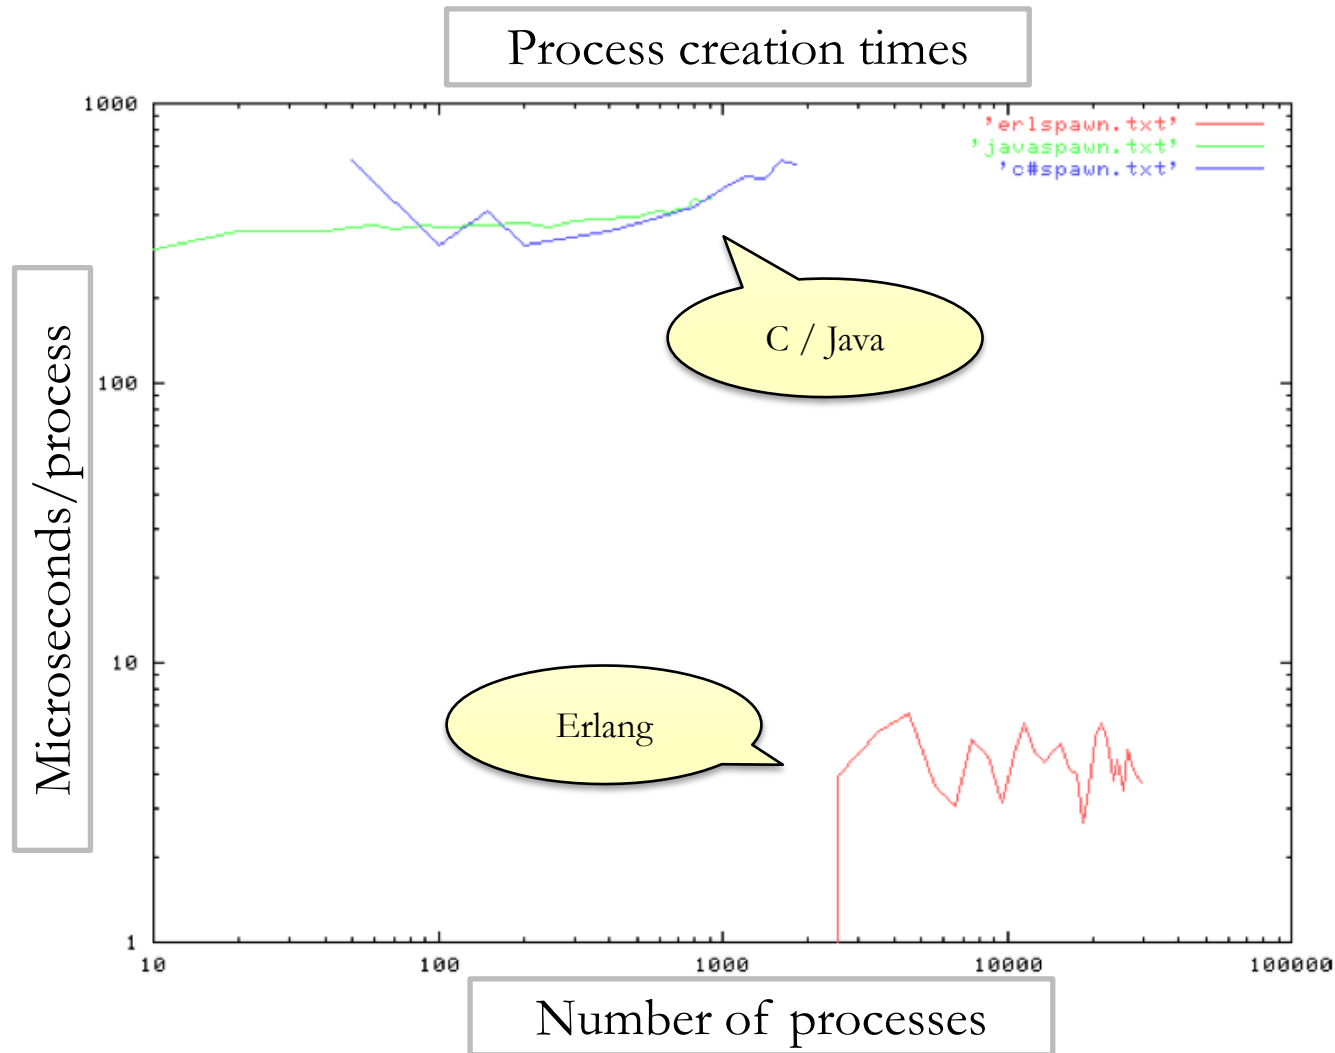


Fault tolerance

- What happens when a process receives an exit signal?
 - If the receiver hasn't declare itself as **system process**, the message will cause it too to exit
 - Otherwise the message will be processed as a normal one
 - `process_flag(trap_exit, true)` turns a process into a system process

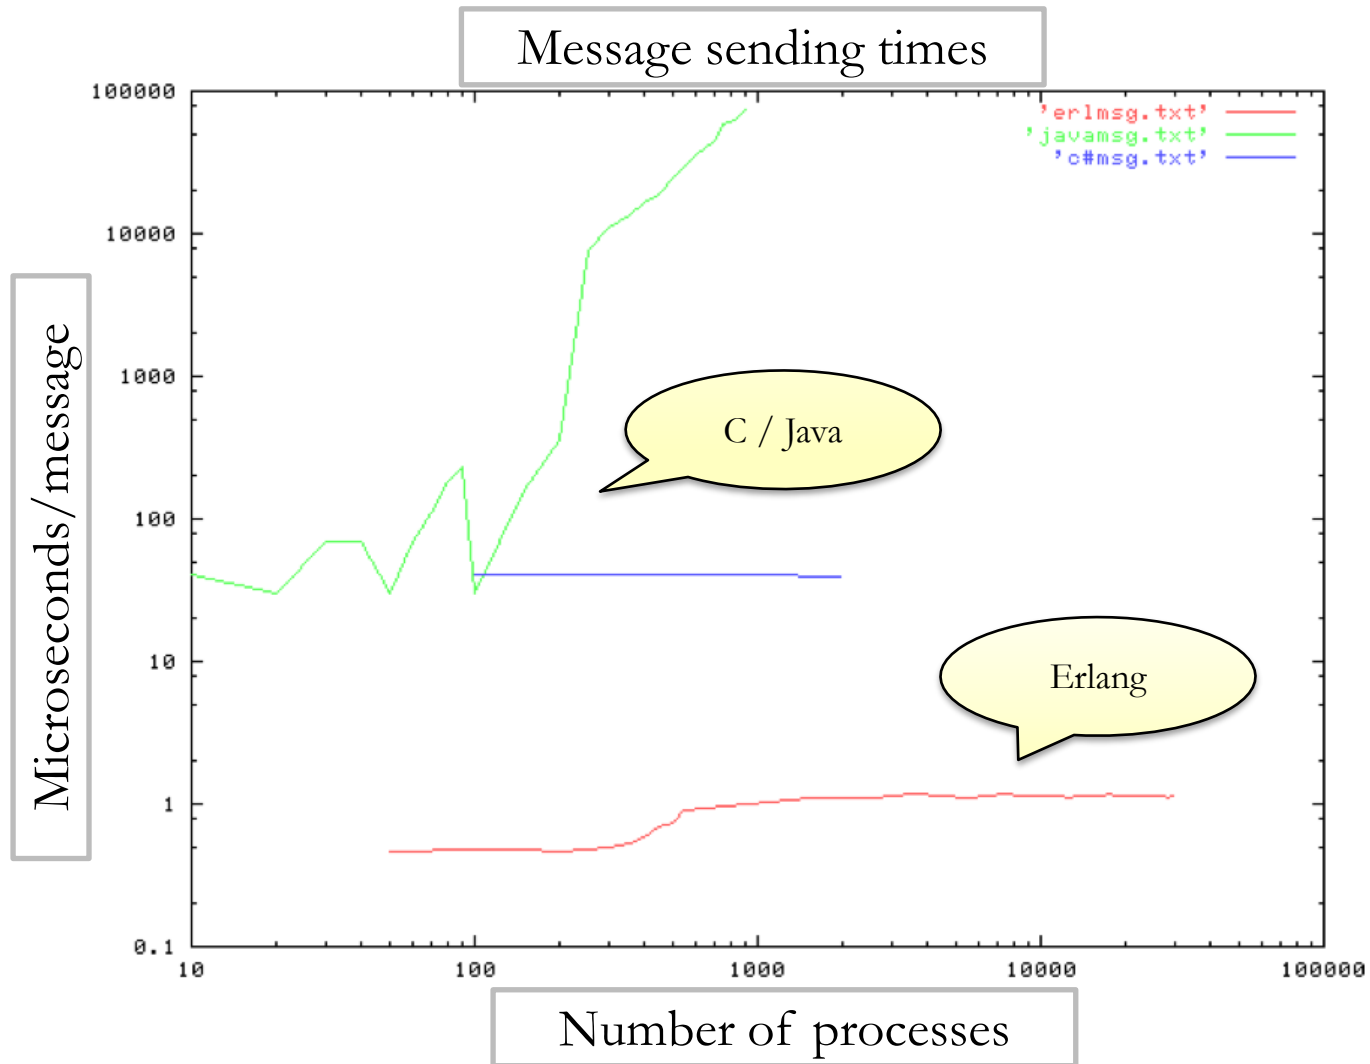


Concurrency: performance



Source: J.Armstrong "Concurrency oriented programming in Erlang"

Concurrency: performance



Source: J.Armstrong “Concurrency oriented programming in Erlang”

Distributed programming

- Erlang enables the programmers to distribute concurrent processes on different machines
- We will talk about two main distribution models:
 - ***Distributed Erlang***: provides a method for programming applications that run on a single administrative domain (trusted environment, like a LAN)
 - Processes run in Erlang ***nodes***
 - All mechanisms for message passing, error handling etc. works as in the single node scenario
 - ***Socket-based distribution***: uses TCP/IP sockets to send messages in an untrusted environment
 - Less powerful but more secure model
 - Used to easily interact with programs written in other languages

Distributed Erlang

- Distributed Erlang enables programmers to spread processes on different *nodes*.
- A node *a* can communicate with a node *b* if
 - It knows *b*'s name
 - *a* and *b* share the same *cookie*
- To start a node with a given name and cookie run
 - `erl -sname name -setcookie cookie (same host)`
 - `erl -name name@host -setcookie cookie (across hosts)`
- We will see how distribution works with two examples:
 - Sending messages to remote nodes
 - Implementing the Remote Procedure Call (RPC) pattern
 - Spawning processes on remote nodes

Distributed Erlang

- Set up your environment
 - Start Erlang with the `-name` option
 - Ensure that both nodes have the same cookie
 - Start the node with `-setcookie cookie`
 - Make sure that the hostnames are resolvable
 - By DNS or
 - Adding an entry to `/etc/hosts`
 - Make sure that all systems have the same version of the code
 - Manually copy the bytecode or
 - Use the shell command `nl(Mod)`
 - It loads the module `Mod` on all connected nodes
 - Useful for dynamic code loading
 - Test if everything is working using the `ping` function
 - `net_adm:ping(node_name)`



Fault tolerant server

- Two processes
 - Server
 - Error handler
- If the server crashes, the error handler
 - Traps the error message
 - Starts a new server
 - Terminates
- Every time a new server starts it creates a new error handler

```
-module(dist).  
-export([start/0, ask/3]).
```

```
start() ->  
    Pid = spawn(fun loop/0),  
    register(server, Pid),  
    io:fwrite("Server started\n"),  
    on_exit(Pid, fun start/0),  
    io:fwrite("Error handler started\n").
```

```
loop() -> receive  
    {Sender, Fun, Num} ->  
        Sender ! Fun(Num),  
        loop()  
end.
```

```
on_exit(Pid, Fun) ->  
    spawn(fun() ->  
        process_flag(trap_exit, true),  
        link(Pid),  
        receive  
            {'EXIT', Pid, _} ->  
                io:fwrite("Error trapped\n"),  
                Fun()  
        end  
    end).  
end).
```



Fault tolerant server

- The client can call the `dist:ask` function to send a message to the server
- If the server does not respond in 100ms, then it returns 'Server crashed'
- Suppose our client wants to execute the function `support:dup`
 - The module has not been loaded by the server

Sends the message to the process
registered as 'server' on the
'Host' machine

```
ask(Host, Fun, Num) ->  
  {server, Host} ! {self(), Fun, Num},  
  receive  
    Reply -> Reply  
  after 100 -> 'Server crashed'  
  end.
```

```
-module(support).  
-export([dup/1]).
```

```
dup(X) -> 2*X.
```




Fault tolerant server

Client's shell

```
>net_adm:ping(server@serverhost).
```

Check the server

```
pong
```

```
>nl(support).
```

Load the module on the server

```
abcast
```

```
>dist:ask('server@serverhost', fun support:dup/1, 16).
```

```
32
```

```
>dist:ask('server@serverhost', fun support:dup/1, aaa).
```

```
'Server crashed'
```

```
>dist:ask('server@serverhost', fun support:dup/1, 24).
```

```
48
```

The server is up
and working again!

Call 2 contains an error.
The server crashes!



Fault tolerant server

Server's shell

```
>dist:start().
```

```
Server started
```

```
Error handler started
```

```
ok
```

Starts the server

```
Error trapped
```

```
Server started
```

```
Error handler started
```

When request 2 is processed,
the wrong message format
crashes the server

The error is trapped and the server
(together with a new error handler)
is started



Distributed Erlang

- Distributed Erlang also enables programmers to spawn processes on a specific node
- Without modifying our previous code we can start the server from the client shell

```
>net_adm:ping(server@serverhost).  
pong  
>nl(dist).  
abcast  
>nl(support).  
abcast  
>spawn(server@serverhost, dist, start, []).  
>dist:ask('server@serverhost', fun support:dup/1, 16).  
32
```

Code mobility

- The `nl` instruction loads a module on a remote machine
 - This realizes code mobility
- It is also possible to move locally defined functions, like in the example below
- This works only on the same host, otherwise the two hosts have to share the bytecode!

```
>Square = fun(X) -> X*X end.
```

Defines a local function

```
>dist:ask('server@serverhost', Square, 10).
```

```
100
```

```
>dist:ask('server@serverhost', Square, aaa).
```

```
'Server crashed'
```


```
>dist:ask('server@serverhost', Square, 12).
```

```
144
```

Uses the function to call the server

A complete example

- Topic based publish-subscribe
 - The dispatcher receives subscriptions and publications from clients
 - It sends published messages to interested (subscribed) clients
 - Similar to topic based communication in JMS





Exported functions

- We define a module that exposes the following functions:
 - startDispatcher: starts the dispatcher of messages, that waits for publish/subscribe commands
 - startClient: starts a process on the client that waits for messages
 - publish: publishes a message for a given topic
 - subscribe: used by clients to express their interests

```
-module(pubsub).  
-export([startDispatcher/0, startClient/0,  
        subscribe/2, publish/3]).
```

```
startClient() ->  
    Pid = spawn(fun clientLoop/0),  
    register(client, Pid).
```

```
clientLoop() ->  
    receive {Topic, Message} ->  
        io:fwrite("Received message ~w for topic ~w~n",  
                  [Message, Topic]),  
        clientLoop()  
    end.
```

```
subscribe(Host, Topic) ->  
    {dispatcher, Host} ! {subscribe, node(), Topic}.
```

```
publish(Host, Topic, Message) ->  
    {dispatcher, Host} ! {publish, Topic, Message}.
```

```
startDispatcher() ->  
    Pid = spawn(fun dispatcherLoop/0),  
    register(dispatcher, Pid).
```



The dispatcher

```
dispatcherLoop() ->  
  io:fwrite("Dispatcher started\n"),  
  dispatcherLoop([]).  
dispatcherLoop(Interests) ->  
  receive  
    {subscribe, Client, Topic} ->  
      dispatcherLoop(addInterest(Interests, Client, Topic));  
    {publish, Topic, Message} ->  
      Destinations = computeDestinations(Topic, Interests),  
      send(Topic, Message, Destinations),  
      dispatcherLoop(Interests)  
  end.
```

- The dispatcher keeps a list of interests
 - Organized as a list of tuples {ClientID, ClientTopicsList}
 - Modified when a subscription is received
 - Used to compute the destinations of a published message



The dispatcher

```
computeDestinations(Topic, Interests) ->  
  computeDestinations(Topic, Interests, []).  
computeDestinations(_, [], Result) -> Result;  
computeDestinations(Topic, [{Client, Interests}|T], Result) ->  
  Matches = matches(Topic, Interests),  
  if Matches == yes ->  
    computeDestinations(Topic, T, Result ++ [Client]);  
  Matches == no ->  
    computeDestinations(Topic, T, Result)  
end.
```

```
matches(_, []) -> no;  
matches(Topic, [H|_]) when Topic == H -> yes;  
matches(Topic, [_|T]) -> matches(Topic, T).
```

```
send(_, _, []) -> ok;  
send(Topic, Message, [Client|T]) ->  
  {client, Client} ! {Topic, Message},  
  send(Topic, Message, T).
```




The dispatcher

Recursively analyze interests
copying them into “Result”

Until “Interests”
becomes empty

```
addInterest(Interests, Client, Topic) ->
  addInterest(Interests, Client, Topic, []).
addInterest([], Client, Topic, Result) ->
  Result ++ [{Client, [Topic]}];
addInterest([{{SelectedClient, Interests}}|T], Client, Topic, Result) ->
  if SelectedClient == Client ->
    NewInterests = Interests ++ [Topic],
    Result ++ [{Client, NewInterests}] ++ T;
  SelectedClient /= Client ->
    addInterest(T, Client, Topic, Result ++ [{SelectedClient, Interests}])
end.
```

Or until the client identifier is found and the
new topic is added to the list of interests
(First if clause)

The dispatcher: improvements

- Is our implementation of the dispatcher efficient?
- We use a list of { *Client*, *Interests* } to represent the interest table
- When we process a publish message we need to check the topic in the *Interests* list of every client
- We can easily modify our code to store, for each topic, the set of interested clients ...
- ... using a list of { *Topic*, *Clients* }
- We only have to slightly modify 3 functions



The dispatcher: improvements

```
computeDestinations(_, []) -> [];  
computeDestinations(Topic, [{SelectedTopic, Clients}|T]) ->  
  if SelectedTopic == Topic -> Clients;  
    SelectedTopic /= Topic -> computeDestinations(Topic, T)  
  end.  
  
send(_, _, []) -> ok;  
send(Topic, Message, [Client|T]) ->  
  {client, Client} ! {Topic, Message},  
  send(Topic, Message, T).  
  
addInterest(Interests, Client, Topic) ->  
  addInterest(Interests, Client, Topic, []).  
addInterest([], Client, Topic, Result) ->  
  Result ++ [{Topic, [Client]}];  
addInterest([{{SelectedTopic, Clients}|T}, Client, Topic, Result) ->  
  if SelectedTopic == Topic ->  
    NewClients = Clients ++ [Client],  
    Result ++ [{Topic, NewClients}] ++ T;  
    SelectedTopic /= Topic ->  
      addInterest(T, Client, Topic, Result ++ [{SelectedTopic, Clients}])  
  end.
```



The dispatcher: now with stdlib

```
dispatcherLoop(Interests) ->
  receive
    {subscribe, Client, Topic} ->
      dispatcherLoop(addInterest(Interests, Client, Topic));
    {publish, Topic, Message} ->
      Destinations = computeDestinations(Topic, Interests),
      send(Topic, Message, Destinations),
      dispatcherLoop(Interests)
  end.

computeDestinations(Topic, Interests) ->
  dict:fold(fun(Client, Current, AccIn) ->
    case lists:member(Topic, Current) of
      true ->
        [Client|AccIn];
      false ->
        AccIn
    end end, [], Interests).

send(Topic, Message, Destinations) ->
  lists:foreach(fun(Client) -> {client, Client} ! {Topic, Message} end, Destinations).

addInterest(Interests, Client, Topic) ->
  dict:update(Client, fun(Current) -> [Topic|Current] end, [Topic], Interests).
```



Socket based distribution

- Erlang offers facility for socket communications
 - We introduce them using a single example (echo server)
 - This enables interaction with other programming languages

```
-module(echo).  
-export([listen/1]).  
  
-define(TCP_OPTIONS,[list, {packet, 0}, {active, false}, {reuseaddr, true}]).  
  
listen(Port) ->  
    {ok, LSocket} = gen_tcp:listen(Port, ?TCP_OPTIONS),  
    {ok, Socket} = gen_tcp:accept(LSocket),  
    do_echo(Socket).  
  
do_echo(Socket) ->  
    case gen_tcp:recv(Socket, 0) of  
        {ok, Data} ->  
            gen_tcp:send(Socket, Data),  
            do_echo(Socket);  
        {error, closed} ->  
            ok  
    end.
```



Socket based distribution

- On the server side:
 - Listen
 - Accept (blocking)
 - Receive
- On the client side:
 - Connect
 - Send
 - Receive

```
> {ok, S} = gen_tcp:connect("localhost", 9000, [{active, false}, {packet, 2}]).
{ok, #Port<0.448>}
> gen_tcp:send(S, "Hello").
ok
> {ok, R} = gen_tcp:recv(S, 0).
{ok, "Hello"}
> R
"Hello"
```



Outline

- Introduction
- Sequential Programming
 - Data structures
 - Single assignment
 - Pattern matching
 - Functional abstractions
 - Dynamic code loading
- Concurrent / Distributed Programming
 - Processes
 - Fault tolerance
 - Distributed Erlang
 - Socket based distribution
- **OTP Introduction**

Open Telecom Platform (OTP)

- What is OTP (Open Telecom Platform)?
 - A set of design principles
 - A set of libraries
 - Developed and used by Ericsson to build large-scale, fault-tolerant, distributed applications
 - It also offers different powerful tools:
 - A complete Web Server
 - An FTP Server
 - A CORBA ORB
 - ...

OTP Behaviors

- There exist structures/patterns used in a great number of different programs
 - Client / Server
 - Server waits for client commands, execute and return responses
 - Worker / Supervisor
 - Workers are processes that perform the computation
 - Supervisors monitor the behavior of workers
 - React when errors are detected (e.g. by restarting the worker)
 - Hierarchies (trees) of supervisors can be created as well
 - Event Manager / Handlers
 - Similar to Java listeners or to publish-subscribe paradigm
 - The manager detects an event
 - The handlers process the event



OTP Behaviors

- Let's take, for example, the client server paradigm
- What varies in different applications adopting this design paradigm?
 - Basically, what the server does
 - The *functional* part of the problem
 - The structure is fixed
 - The *non-functional* part of the problem
- The idea is to use higher order functions abstraction
 - The common non-functional part is implemented in modules called *behaviors*
 - The functional part has to be implemented in modules that export predefined functions
 - *Callback functions*
- Do not reinvent the wheel!

Applications

- OTP dictates also a common structure for *applications* i.e. pieces of code providing a specific functionality
- Following this structure applications can be:
 - Started, stopped, configured and monitored as a unit
 - Reused to build higher level applications
 - Included applications
- Often applications are defined as distributed
 - Run on different cooperating nodes
 - Realize fault tolerance using distributed worker/supervisor pattern
- This simplifies the design of component based architectures where different functional units can be combined to solve a complex task



Release handling

- Applications come with a release resource file that defines dependencies between applications
- It is possible to express dependencies involving the versions of considered applications
- Release handling tools
 - Start from release resource files
 - Can generate automatic procedures to update a particular application
 - Automatic resolution of dependencies
 - Based on low level dynamic code loading
 - Work in distributed scenario
 - Try to upgrade without stopping involved applications
 - Not always possible
 - Sometimes it is necessary to restart the application after upgrade
 - Not always easy to configure correctly



References

- J. Armstrong: “*Programming Erlang: Software for a Concurrent World*”. Pragmatic bookshelf, 2007.
- R. Virding, C. Wikstrom, M. Williams: “*Concurrent Programming in Erlang (2nd Ed)*”. Prentice Hall Intl, 1996.
- J. Armstrong: “*A History of Erlang*”. In Proceedings of the 3rd conference on History of Programming Languages, 2007.
- J. Armstrong: “*Concurrency Oriented Programming in Erlang*”. Invited paper in FFG, 2003.
- C. Wikstrom: “*Distributed Programming in Erlang*”. In Proceedings of the 1st International Symposium on Parallel Symbolic Computation, 1994.
- J. Larson: “*Erlang for Concurrent Programming*”. Queue 2008.
- Documentation available at:
 - <http://www.erlang.org/doc.html>
- OTP Team “Design Principles”