



# Distributed Systems Security

**Alessandro Sivieri**

Dipartimento di Elettronica e Informazione  
Politecnico di Milano, Italy

[sivieri@elet.polimi.it](mailto:sivieri@elet.polimi.it)

<http://corsi.dei.polimi.it/distsys>

Slides based on previous works by Matteo Migliavacca and  
Alessandro Margara



# Outline

- Introduction
- Cryptography
- Secure Channels
- Secure Group Communication
- Security Management
- Access Control
- Electronic Payment



# Introduction



# Introduction

- Why security in a distributed systems course?
- Sharing of resources is the motivating factor for distributed systems
- Processes encapsulate resources and provide access to them through interaction with other processes
- We want this interaction to be “correct”
  - Protect resources against unauthorized accesses
  - Secure processes against unauthorized interactions
  - Secure communication (messages between processes)



# Introduction

- We postulate an enemy (adversary) capable of:
  - Sending any message to any process
  - Reading and copying any message between a pair of processes
- Potentially any computer connected to the network
  - In an authorized way
  - In an unauthorized way
    - Through a stolen account
- Security threats from the enemy
  - Interception (sniffing, dumping)
  - Interruption (disruption, denial of service)
  - Modification (tampering, website defacing)
  - Fabrication (injection, replay attacks)



# Introduction

- Security to ensure that the overall system meets these requirements
  - Availability
    - We want the system ready for use as soon as we need it
  - Reliability
    - The system should be able to run continuously for long time
  - Safety
    - The system should cause nothing bad to happen
  - Maintainability
    - The system should be easy to fix
  - Confidentiality
    - Information should be disclosed to authorized party only
  - Integrity
    - Alteration can be made only in authorized ways
      - data alteration but also hardware and software



# Introduction

- We said *correct* interaction, *authorized* access ...
- ... but *correct* and *authorized* wrt what?
- We need to specify a Security Policy
  - Defines precisely what is allowed and what's forbidden



# Introduction

A security policy can be enforced through

- Security Mechanisms
  - Encryption (confidentiality, verify integrity)
  - Authentication (identification of parties)
  - Authorization (control information access)
  - Auditing (breach analysis, but also IDS)

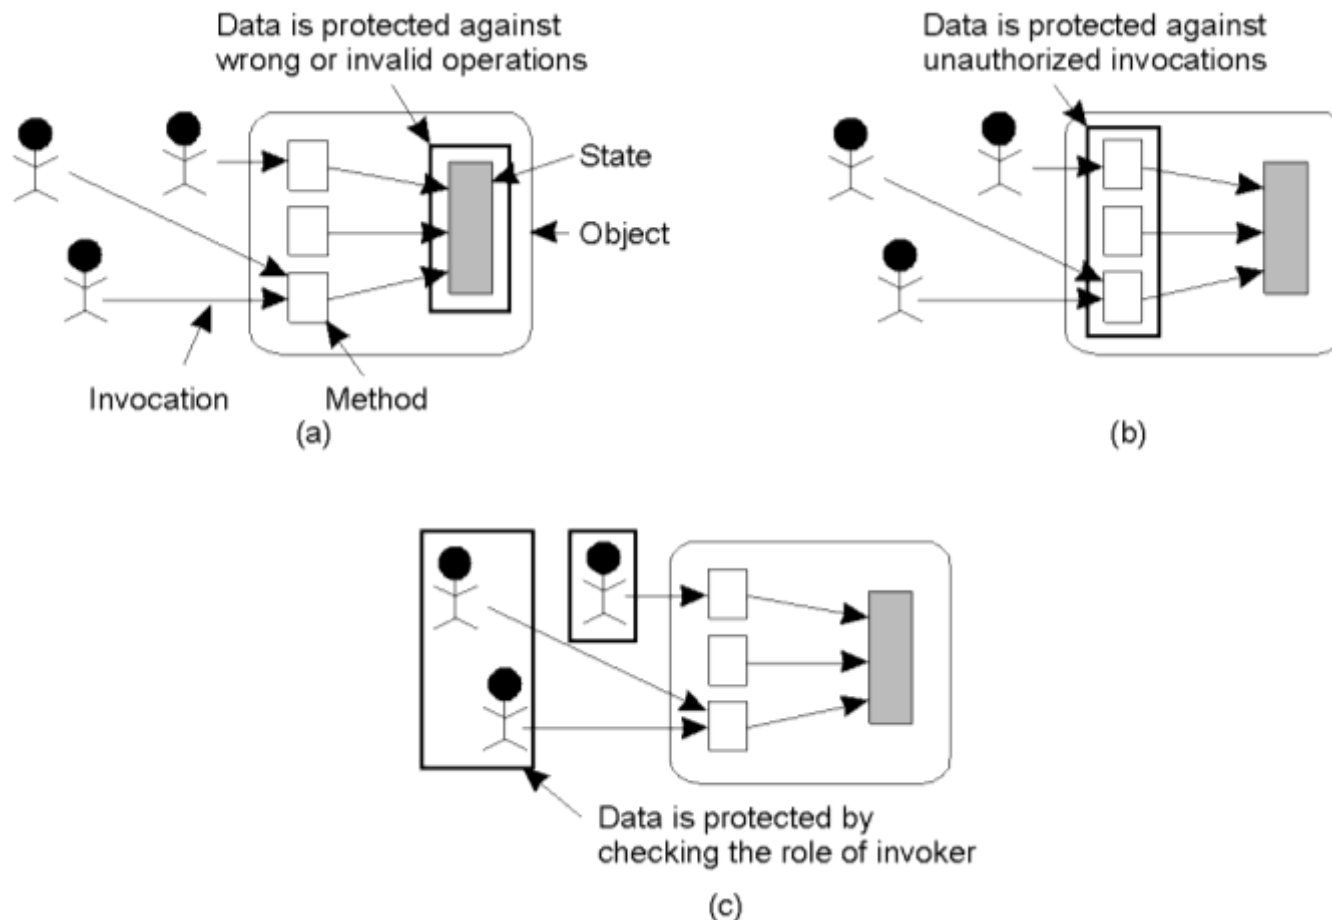


# Design Issues

- When considering the protection of a (distributed) application, there are essentially three approaches that can be followed:
  - Concentrate directly on protection of the data, irrespective of the various operations that can possibly be performed on data items
  - Concentrate on protection by specifying exactly which operations may be invoked
  - Focus directly on users: only specific people have access to the application (role of the invoker)

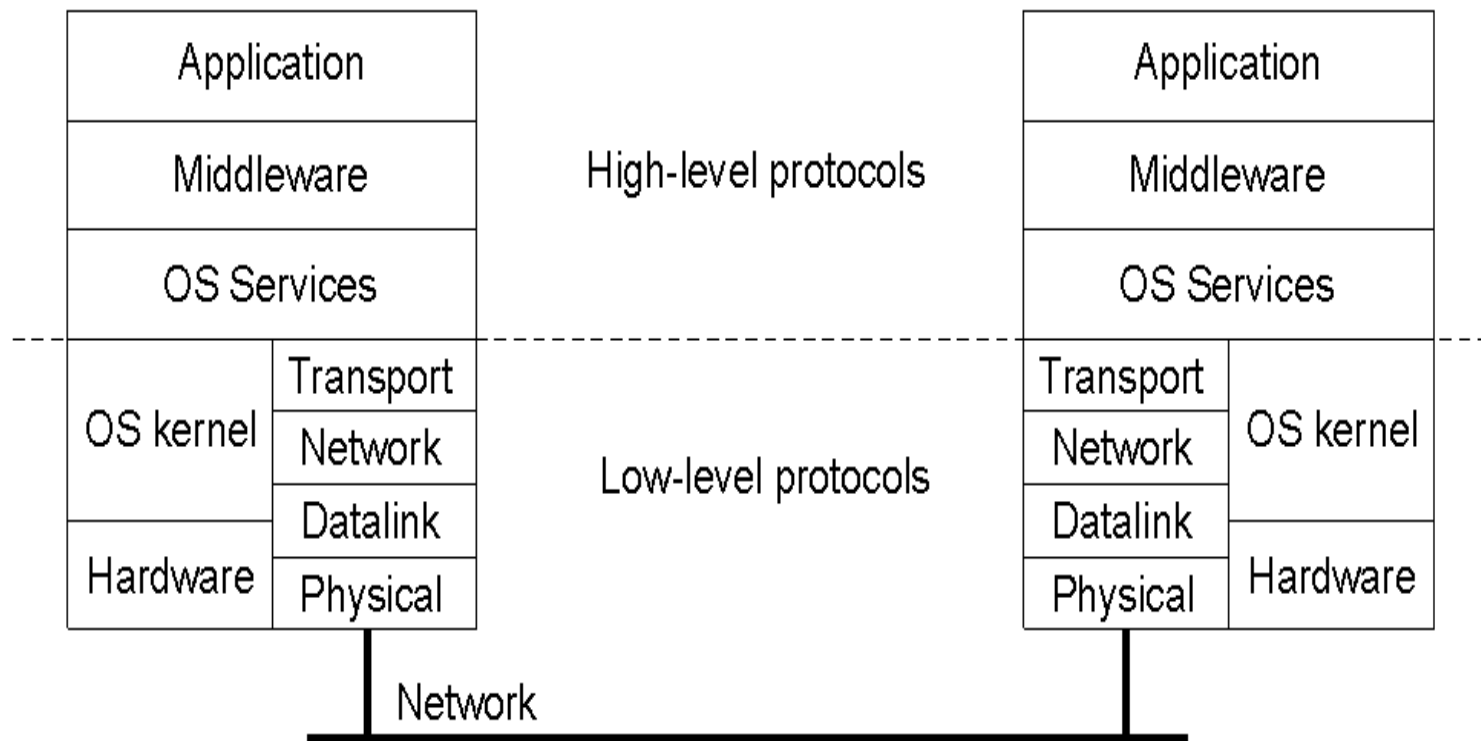
# Design Issues

- Focus of control



# Design Issues

- Layering of security mechanisms

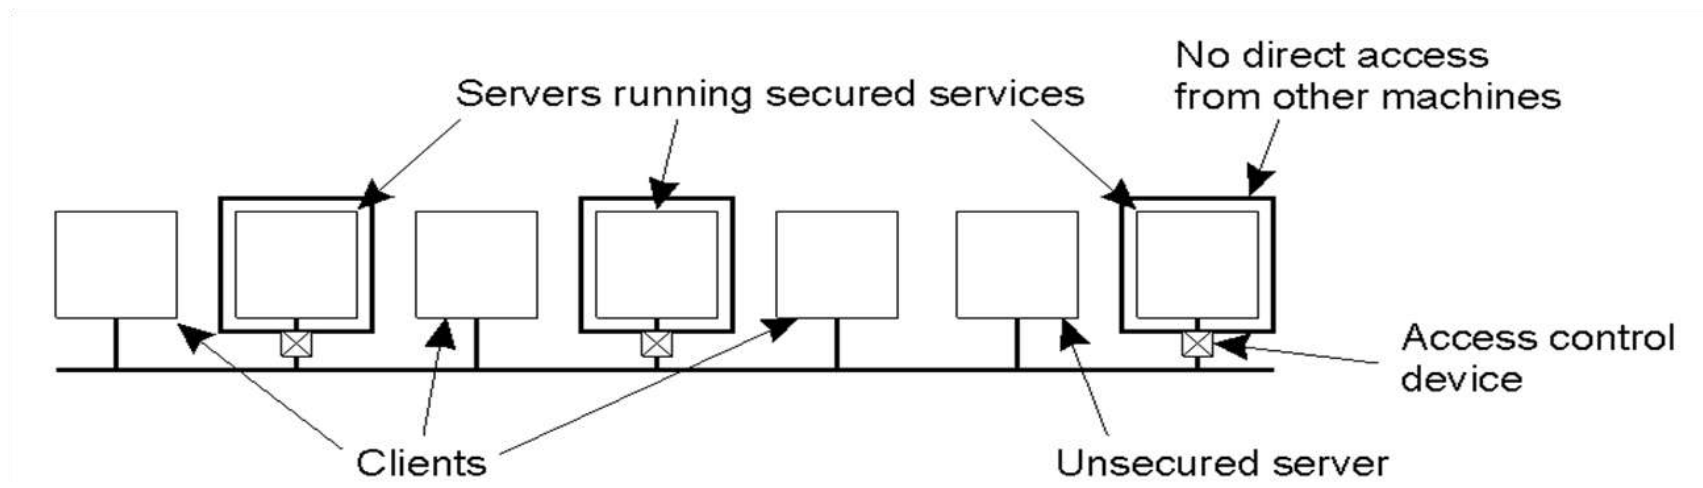


# Layering of security mechanisms

- Where to put security mechanisms?
  - If you don't trust security at low level you can build security mechanisms at higher level
    - If you don't trust transport and lower levels you can use SSL
- What to trust?
  - High level mechanisms depend on lower level mechanisms
    - Netscape's SSL impl was built on top of 40 bit RC4, and hence was easily breakable
    - You need trust in local operating system (kernel at min)
  - The set of mechanisms you need to trust to enforce a given policy is called Trusted Computing Base (TCB)

# TCB

- Usually TCB is composed of many systems
- A way to reduce the TCB is to separate trusted and untrusted services, and granting access to trusted services by a minimal reduced secure interface
  - RISSC: Reduced Interfaces for Secure System Components



# Design Issues

- Simplicity
  - As usual simple mechanisms lead to fewer errors (at design and implementation level)
  - But unfortunately (as we will see in the following) simple mechanisms are often not sufficient to build secure systems
    - Examples of this complexity also emerge at the user level:
      - firewall-related problems
      - certificate verification
      - ...
  - Often applications are complex themselves and security makes it worse



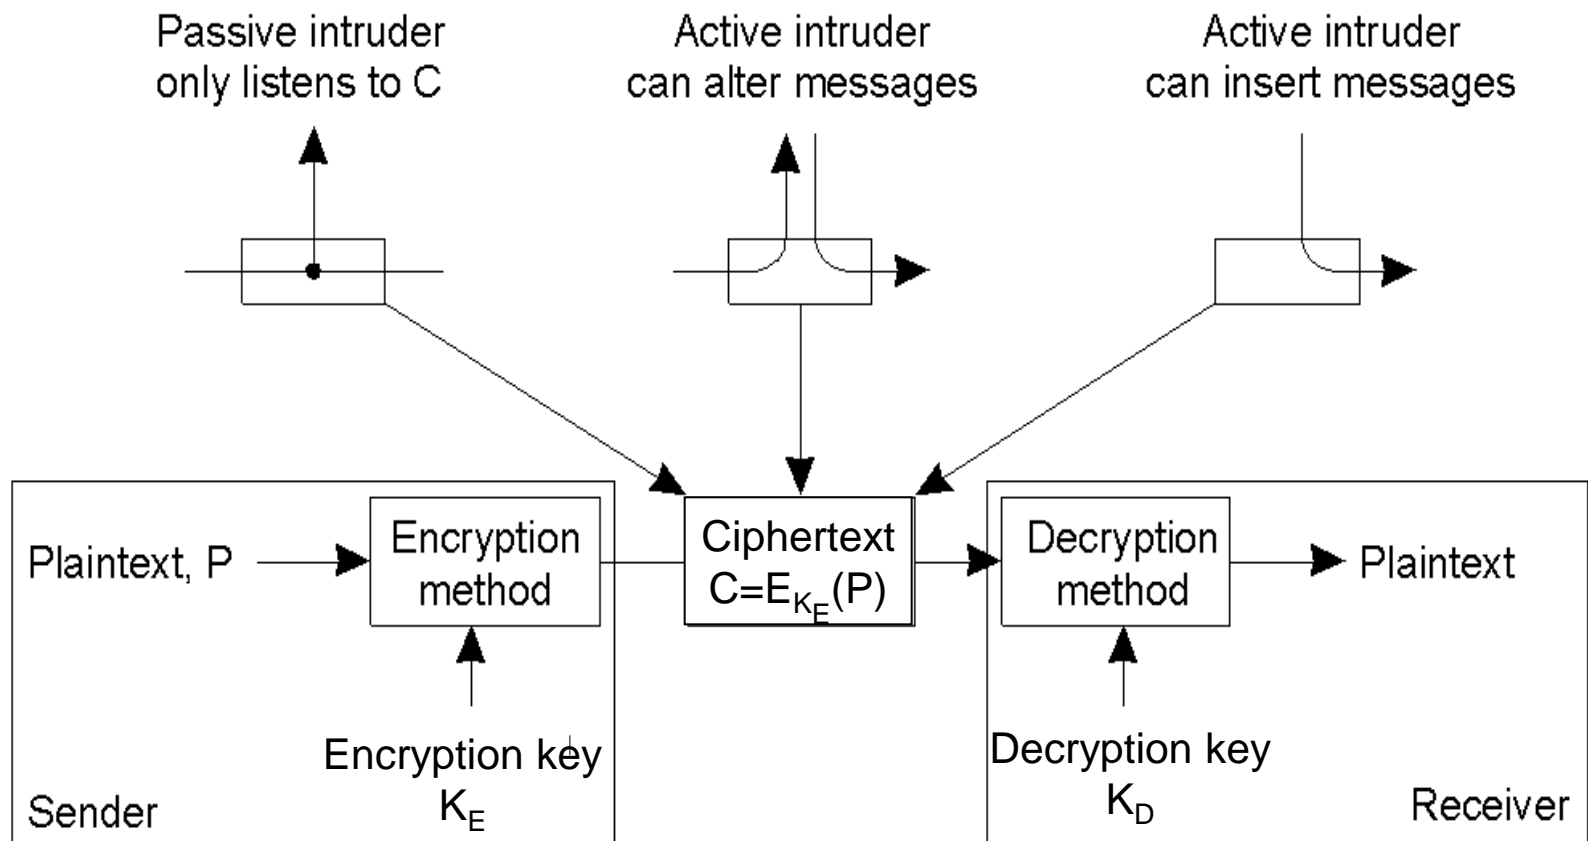
# Cryptography

# Encryption

- Message P – encryption function  $\rightarrow$  Message C
    - P is called Plaintext, C is called Ciphertext
    - The encryption function must be invertible (decryption)
  - Encryption function used to be secret (*security through obscurity*), but this proved to be dangerous
    - Function can't be subject to public review
      - you must trust the inventor of the function
  - Parametric encryption functions
    - From secret functions to known functions with secret parameters (keys)
- $$C = E_K(P) \qquad P = D_K(C) \qquad D_K = E_K^{-1}$$



# Encryption



# Symmetric encryption

- If  $K_E = K_D$
- We use the symbol  $K_{A,B}$  to denote the key shared by A and B
- To achieve communication between N parties we need  $O(N^2)$  different keys

# Asymmetric Encryption

- $K_E \neq K_D$
- $K_E$  and  $K_D$  are uniquely tied to each other (form a pair)
  - $K_D$  can decrypt only from  $K_E$ ,  $K_E$  can encrypt only for  $K_D$
  - Computing  $K_E$  from  $K_D$  or vice-versa must be computationally infeasible
  - We can distribute one key without danger for the other
- So we can safely distribute  $K_E$  to everyone who is interested in sending messages to me (while keeping  $K_D$  private)

# Asymmetric Encryption (cont'd)

- We can also reverse things: distribute  $K_D$  and keep private  $K_E$ 
  - This is how e-signing is done
- So there's no point in calling the two keys of a pair  $K_E$  and  $K_D$ .
- We call them  $K_A^+$  (A's public key) and  $K_A^-$  (A's private key)



# Hash

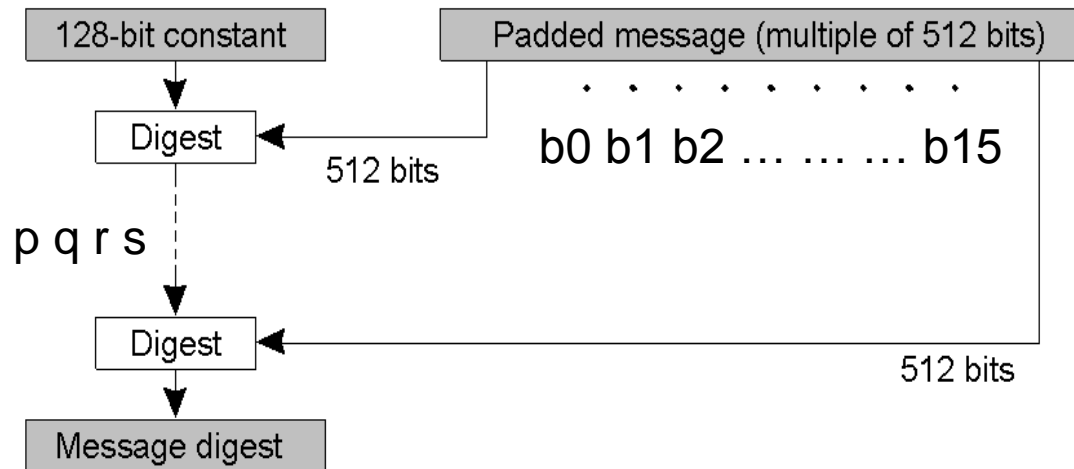
- Message  $m$  – hash function  $\rightarrow$  Hash  $h$ 
  - $h = H(m)$
- A note about hashing (from a math point of view)
  - Since (usually)  $H$  function co-domain (range) is smaller than the domain,  $H$  cannot be injective
  - Collisions:  $m \neq m'$  but  $H(m) = H(m')$ 
    - MD5 uses 128 bit (16 byte). Messages can be of any length  $\Rightarrow$  if we try every 16 byte long message we exhausted the available space so for 17 bytes we have 256 collision, for 1Mbyte files we have 250 Millions collision
  - It's NOT true that for every message there's a different hash!!
  - So why we use hash for data integrity and digital signature?



# Hash (cont'd)

- Message  $m$  – hash function  $\rightarrow$  Hash  $h$ 
  - $h = H(m)$
- Properties
  - Depends on hashing intended use (error codes, hash tables)
  - One way: given  $h$  such that  $h = H(m)$ 
    - it should be hard to find  $m$
  - Weak collision resistance: given  $h$  and  $m$  such that  $h = H(m)$ 
    - it should be hard to find  $m' \neq m$  such that  $h = H(m')$
  - Strong collision resistance: given  $H()$ 
    - it should be hard to find  $m' \neq m$  such that  $H(m) = H(m')$

# MD5 (Message Digest / Hash)



Iterations 1-8	Iterations 9-16
$p \leftarrow (p + F(q, r, s) + b_0 + C_1) \ll 7$	$p \leftarrow (p + F(q, r, s) + b_8 + C_9) \ll 7$
$s \leftarrow (s + F(p, q, r) + b_1 + C_2) \ll 12$	$s \leftarrow (s + F(p, q, r) + b_9 + C_{10}) \ll 12$
$r \leftarrow (r + F(s, p, q) + b_2 + C_3) \ll 17$	$r \leftarrow (r + F(s, p, q) + b_{10} + C_{11}) \ll 17$
$q \leftarrow (q + F(r, s, p) + b_3 + C_4) \ll 22$	$q \leftarrow (q + F(r, s, p) + b_{11} + C_{12}) \ll 22$
$p \leftarrow (p + F(q, r, s) + b_4 + C_5) \ll 7$	$p \leftarrow (p + F(q, r, s) + b_{12} + C_{13}) \ll 7$
$s \leftarrow (s + F(p, q, r) + b_5 + C_6) \ll 12$	$s \leftarrow (s + F(p, q, r) + b_{13} + C_{14}) \ll 12$
$r \leftarrow (r + F(s, p, q) + b_6 + C_7) \ll 17$	$r \leftarrow (r + F(s, p, q) + b_{14} + C_{15}) \ll 17$
$q \leftarrow (q + F(r, s, p) + b_7 + C_8) \ll 22$	$q \leftarrow (q + F(r, s, p) + b_{15} + C_{16}) \ll 22$

Total of 64 iterations for each message part  
Uses 4 functions (F, G, H, I)

# MD5 strong collision resistance

Jun 15, 2004 - Obsolete and Popular Banking Algorithm Gains New Foe

Ottawa - CertainKey, Inc., as a supporter and implementer of strong encryption, is proud to announce that it will award \$10,000 to the first person or group to find a collision\* in MD5, a 128-bit message digest (security algorithm) used by companies such as Bank of America, Citibank, Fleet Bank, eTrade.com USA, and eBay.com.

According to cryptography industry experts, the algorithm became obsolete in 1998, yet remains in use and leaves secure systems vulnerable to attack. With this award, CertainKey hopes to accelerate the adoption of more cryptographically secure algorithms to replace MD5 in all industry applications that relate to cryptography, and encourage public awareness of data security in general.



# MD5 strong collision resistance

<a href="#">Products</a>	<a href="#">Services</a>	<a href="#">Support</a>	<a href="#">About / Contact</a>
--------------------------	--------------------------	-------------------------	---------------------------------



CertainKey  
Cryptosystems

Monday 24th of January 2005

## Official Rules

### Final Results

The challenge is now over.  
A Collision in the full MD5 hash function was found by researchers in China. [Wang, Feng, Lai, and Yu](#) published full collisions for:

- MD4,
- MD5,
- HAVAL-128, and
- RIPEMD-128

This represents one the the most substantial developments in cryptography in recent years. Using their technique, an MD5 collision can be found in a matter or hours!!!  
Here are the files you can use to verify their discovery: [md5col.zip](#)

### Award

CertainKey will give one award of \$10,000 CAD in a lump sum to the individual or group that discovers the first collision in the MD5 hash algorithm as defined by RFC 1321.

Only one award will be given. CertainKey will verify the validity of all submitted claims and the award will be paid for the first verifiable collision.

### GreyWire

Replace T1/T3 trunk lines with a long distance, high-speed, secure wireless link.

### \$10,000 Prize


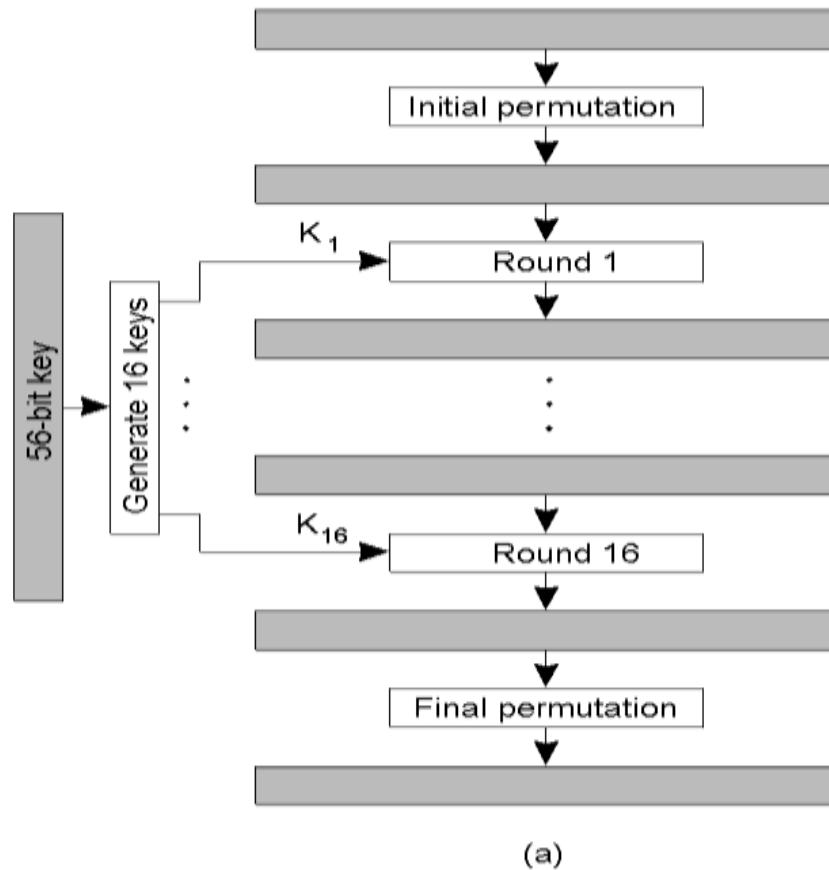
\$10,000 prize claimed for the first MD5 collision.

### Secure Archives

We are offering our Valise™ software free of charge. Upgrades are available to create secure archives as well. Click [here](#) for more information.

# DES (Symmetric key)

Data Encryption Standard. Standard from 1977 to 2001  
Replaced by Advanced Encryption Standard (AES)






# RSA (Public key)

- Based on the difficulty of factorization of large numbers
- Key generation
  - Choose two very large prime numbers,  $p$  and  $q$
  - Compute  $n = p \times q$  and  $z = (p - 1) \times (q - 1)$
  - Choose a number  $d$  that is relatively prime to  $z$
  - Compute the number  $e$  such that  $e \times d = 1 \pmod{z}$  (easy)
- $K_B^- = (d, n)$  and  $K_B^+ = (e, n)$
- Alice encrypts  $m$ :  $c = m^e \pmod{n}$  and transmit  $c$
- Bob decrypts  $c$ :  $m = c^d \pmod{n}$ 
  - It can be proven that  $m = m^{de} \pmod{n}$
- Symmetric key algorithm are faster: RSA 100-1000 times slower than DES

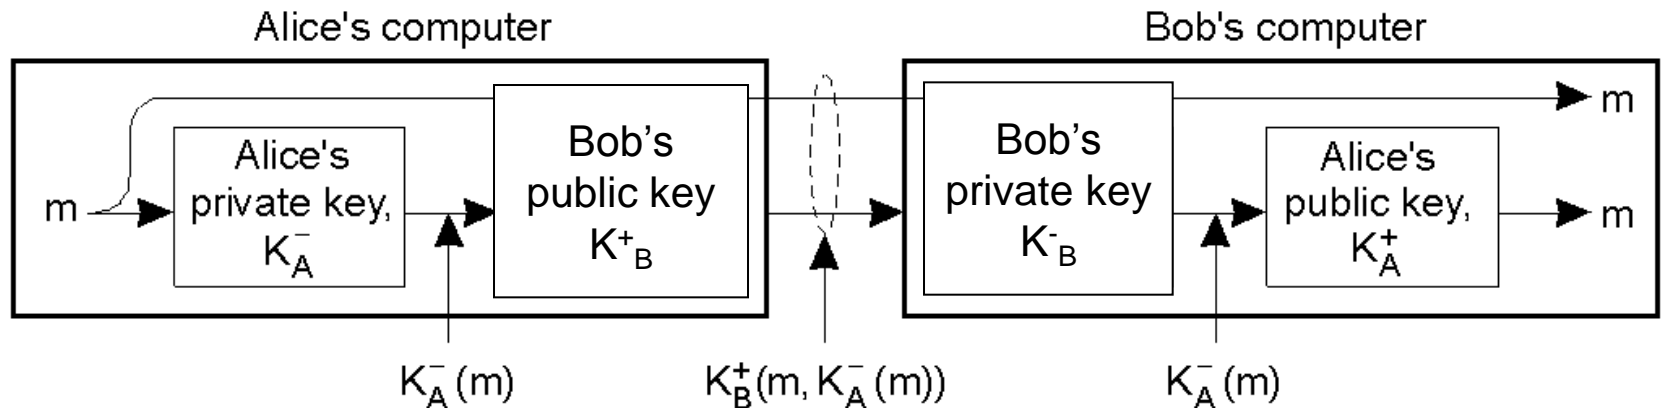
# RSA (Public key) - Example

- Simple example with small prime numbers
  - $p = 11$
  - $q = 5$
  - $n = p \times q = 55$
  - $z = (p-1) \times (q-1) = 40$
  - Choose a number  $d$  that is relatively prime to  $z$ :  $d = 13$
  - Compute the number  $e$  such that  $(e \times d) = 1 \pmod{z}$ :  $e = 37$
- Alice wants to encrypt message  $m = 16$  to send it to Bob
  - Alice's part of the secret is  $(e, n) = (37, 55)$
  - Alice encrypts  $m$ :  $c = m^e \pmod{n} = 36$
  - Alice sends  $c = 36$  to Bob
  - Bob's part of the secret is  $(d, n) = (13, 55)$
  - Bob decrypts  $m$ :  $m = c^d \pmod{n} = 16$

# Digital Signatures

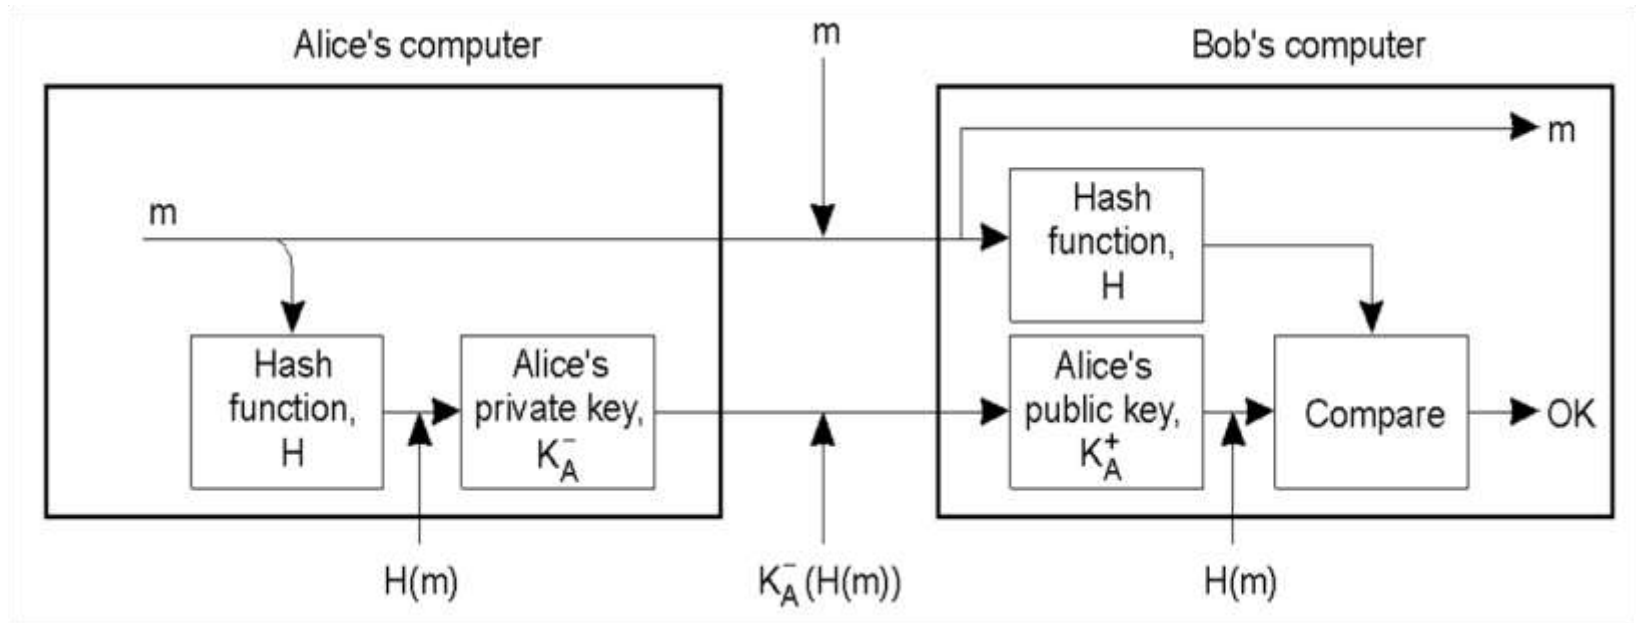


Signing



Signing and encrypting

# Digital Signatures



Digitally signing a message using a message digest

# Certification Authorities

- How can we trust the association *public key* – *physical person*?
- Authentication of public keys
  - This is done by means of public-key certificates
  - A public key certificate is a tuple
    - Identity
    - Public key
    - Signed by a Certification Authority (CA)
  - The public key of CA is assumed to be well-known
    - The basic idea: pervasive information is hard to alter, if a CA public key is everywhere it is really hard to alter every copy without being noticed
  - Yet CA needs to authenticate public keys before issuing a certificate

# Certification Authorities

- We can have several trust models
  - Hierarchical
    - Root CA belongs to central authority (possibly governs)
    - There is a hierarchy of CAs that certificate each other
    - Leaf CAs certificate users
  - PGP's web of trust
    - Users can authenticate other users by signing their public key with their own
    - Users can define who they trust to authenticate others
    - The two are orthogonal: I can be sure that  $K^+A$  is Alice's key but I may not trust her diligence in signing others' key



# Certification Authorities

- Certificates lifetimes
  - Protect from compromise
    - A certificate associates public-key and the owner
- Secret key is compromised. Now what?
  - Revocation
    - Certificate revocation list (CRL) published periodically by the CA
    - Every time a certificate is verified the current CRL should be consulted
    - This means that a client should contact CA every time a new CRL is published



# Secure channels



# Secure channels


- For secure communication in a distributed systems we can exchange information through secure channels
- Secure channels provide protection against
  - Interception (through encryption)
  - Modification and Fabrication (through authentication and message integrity)
  - Do not protect against interruption
- We assume that processes are secure, whereas every message can be intercepted, modified and forged by an attacker



# Secure channels

- Authentication and message integrity should go together (sender and content are together)
  - If a message is modified knowing the sender of the original message is no longer useful
  - If a message is never modified but I do not know the source it's not very useful
- Authentication needs shared information between the authenticator and the party
  - the very concept of authentication requires that (not the implementation in a specific protocol)
  - in the following protocols this information is the authorization key (either symmetric or asymmetric) exchanged beforehand
  - How this authorization key is exchanged? Hard we'll try to answer later
  - Auth protocols tackle how to verify this common information without disclosing it on the channel

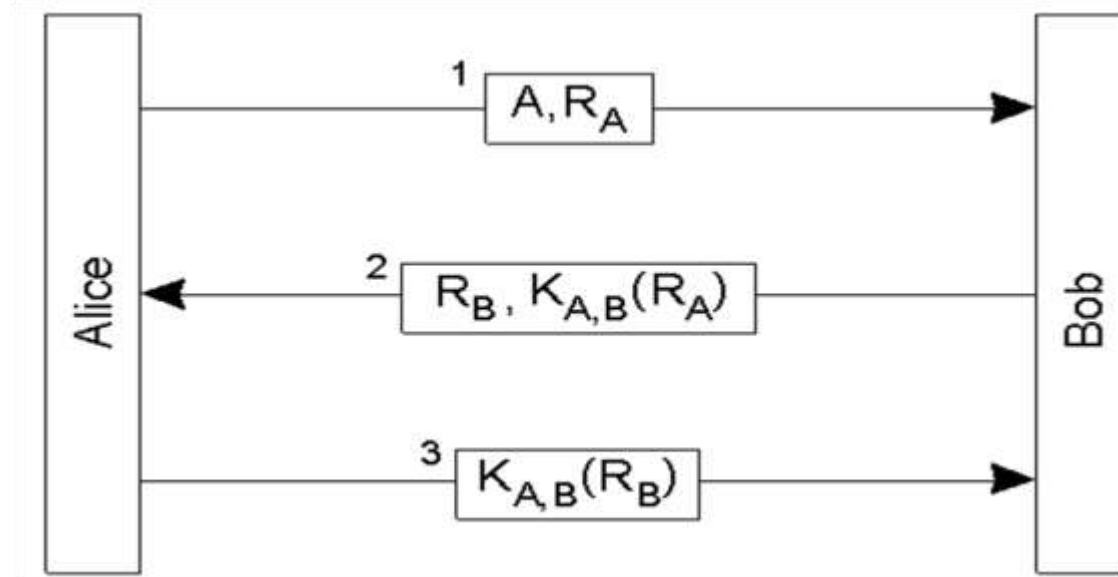
# Secure channels



- Authentication with shared secret key
  - challenge-response



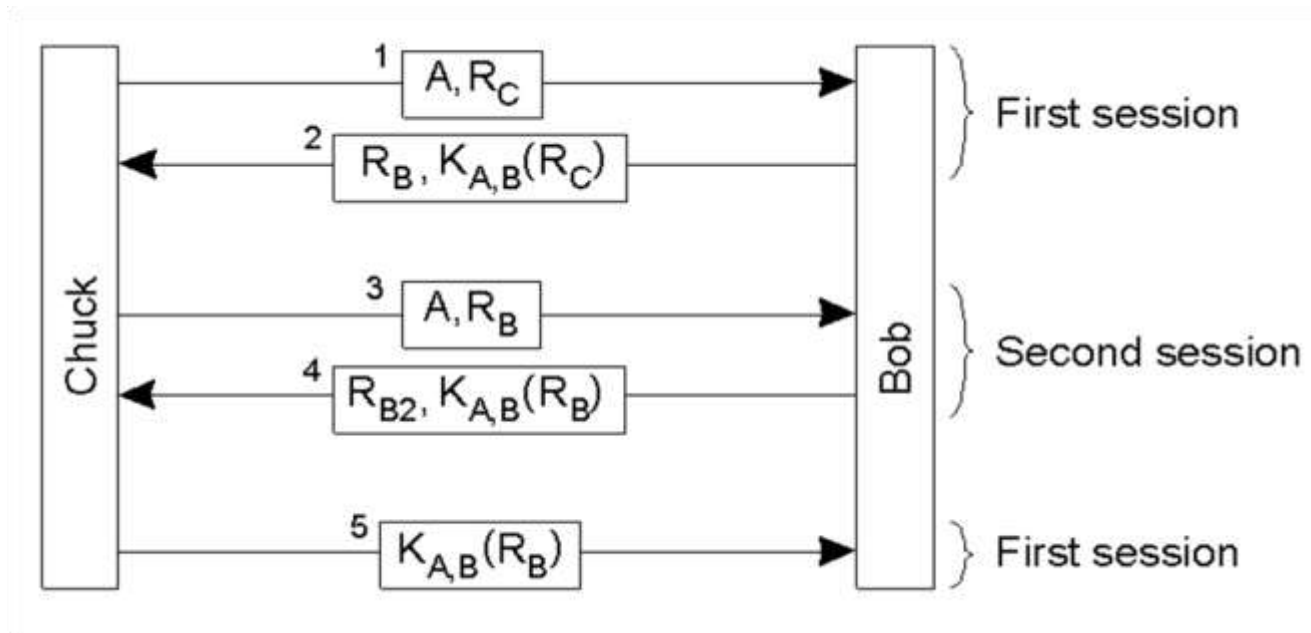
# Secure channels



- It seems a very long exchange... let's try to piggyback some message
- Since A is going to authenticate B, let's challenge B in the first round
- But it does not work...



# Secure channels



- Reflection attack
  - An attacker can request the response of the challenge!
    - Fix: use even for request, odd for response
- Key wearing
  - An active attacker can solicit use of the key (causing “key wearing”)



# Secure channels

- What happens in a secure channel set-up after authentication?
- Usually a session key (symmetric) is exchanged to provide integrity and possibly confidence of following messages
- Session keys are useful to limit the wearing of the main key (used for authentication)
- After session is closed the session key must be destroyed

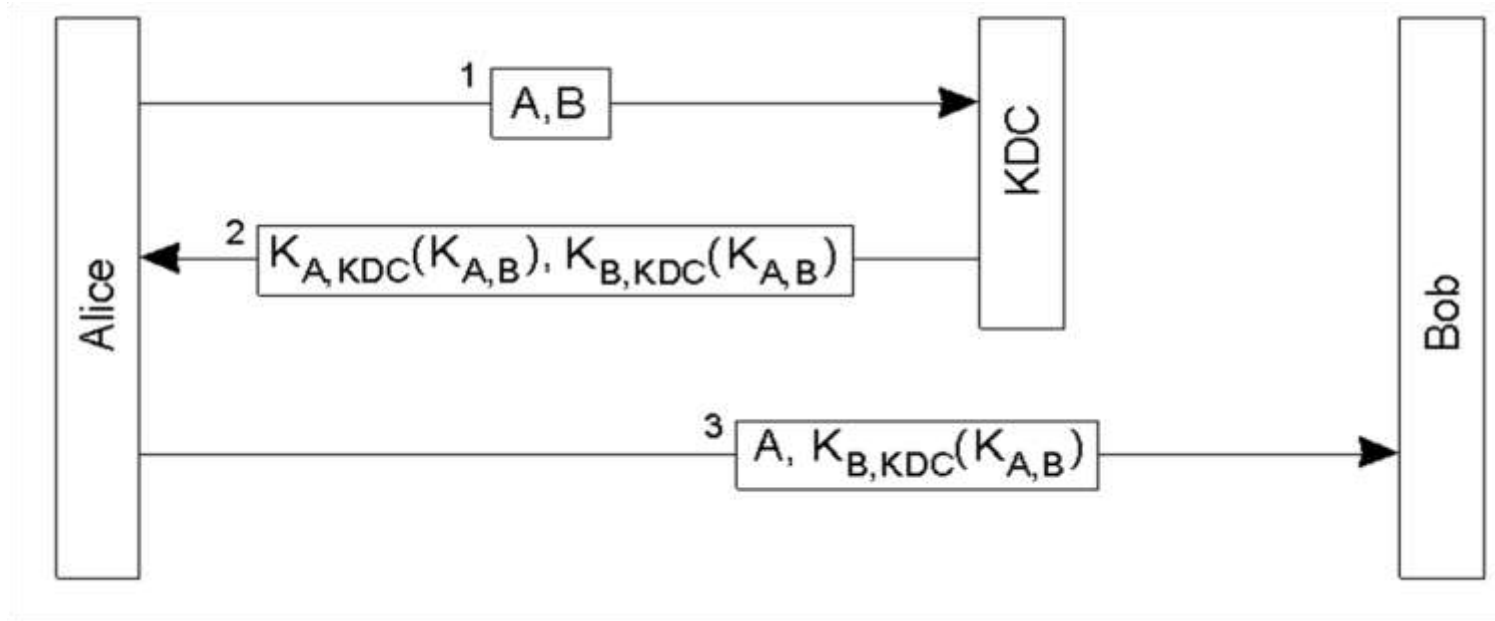


# Secure channels

- One of the problems with using shared secret key for authentication is scalability
  - If a distributed system contains  $N$  hosts, and each host is required to share a secret with each of the other  $N-1$  hosts, we have  $N(N-1)/2$  keys!
  - An alternative is to use a centralized approach by means of a ***Key Distribution Center (KDC)***
  - KDC shares a secret with each of the hosts ( $N$  keys) and generates tickets to allow communication between hosts

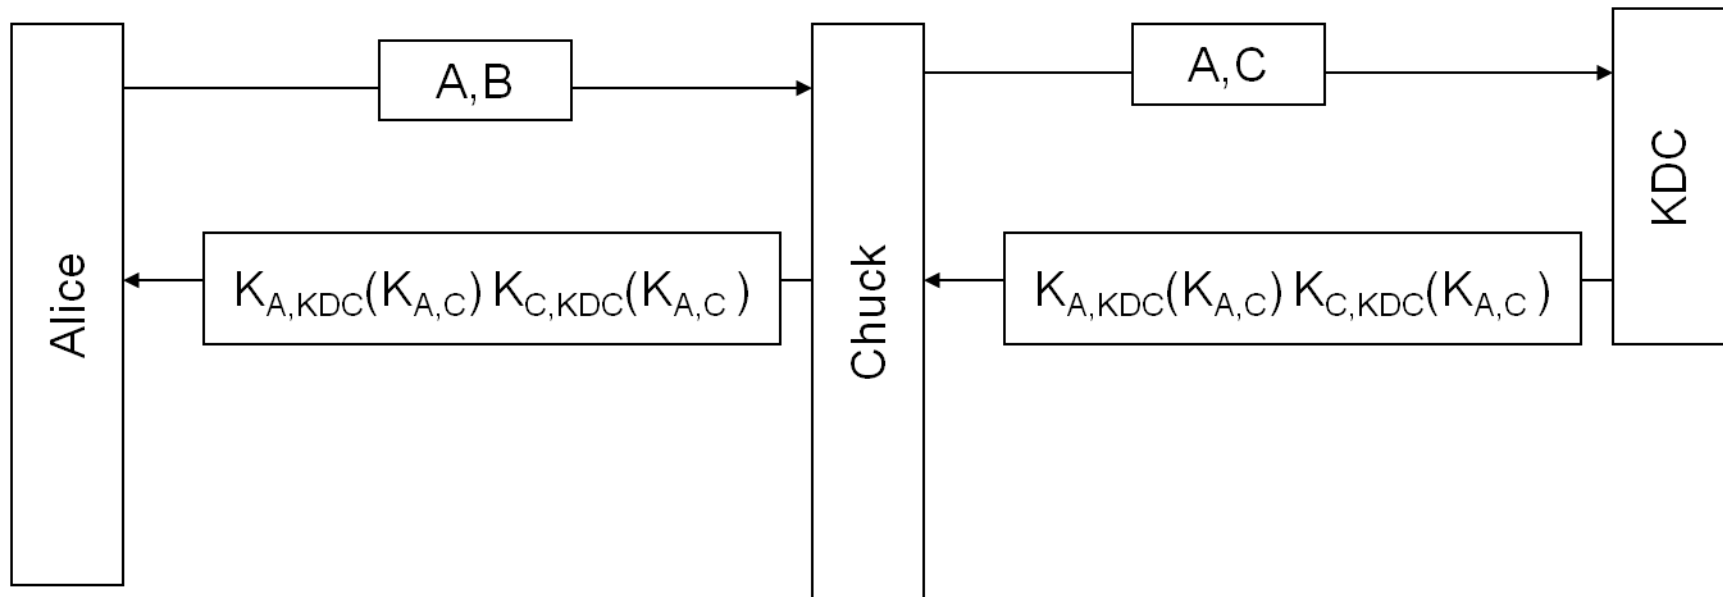


# Secure channels



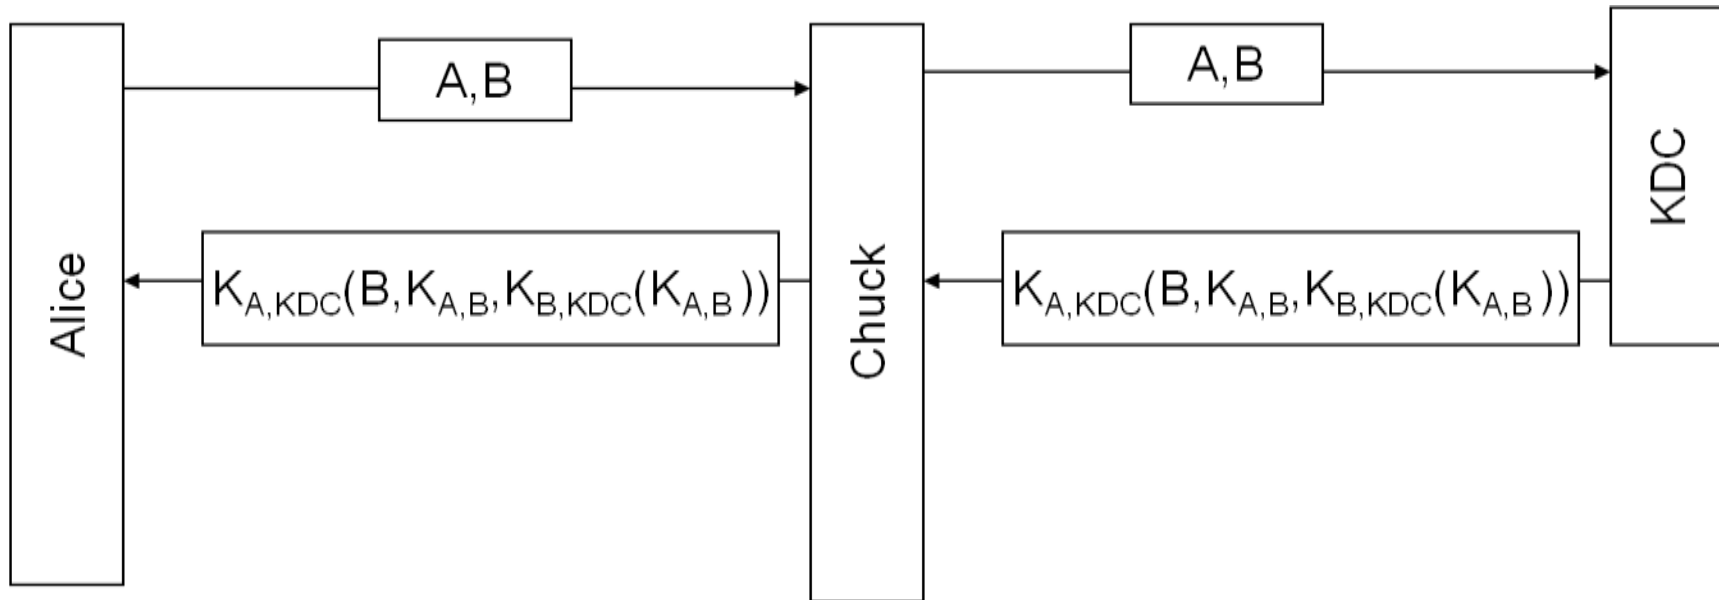
- Auth with trusted KDC
- $K_{B,KDC}(K_{A,B})$  is called ticket
- This protocol is not safe from many attacks

# Secure channels



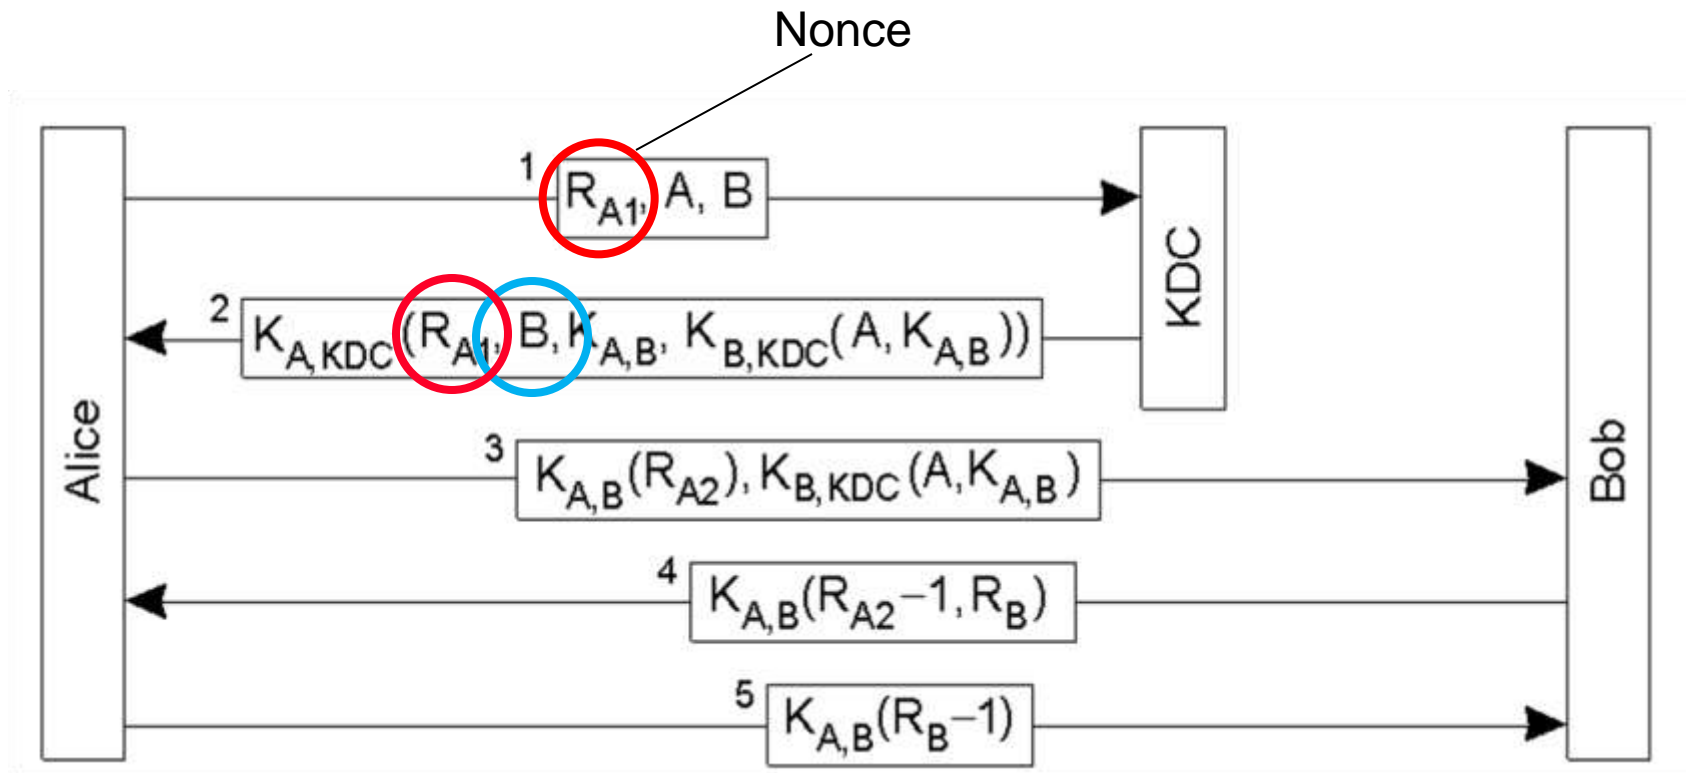
- Modification of the first message
- We could encrypt  $A, B$ :  $K_{A,KDC}(A, B)$
- Unfortunately Chuck can replay an old  $K_{A,KDC}(A, C)$

# Secure channels



- Solution: put  $B$  in the response from KDC and protect it with  $K_{A,KDC}$
- Chuck cannot modify  $A$ 's request.
- But... after stealing  $K_{B,KDC}$  he can serve an old KDC reply forever (even after a fresh negotiation of  $K_{B,KDC}$  between Bob and KDC)

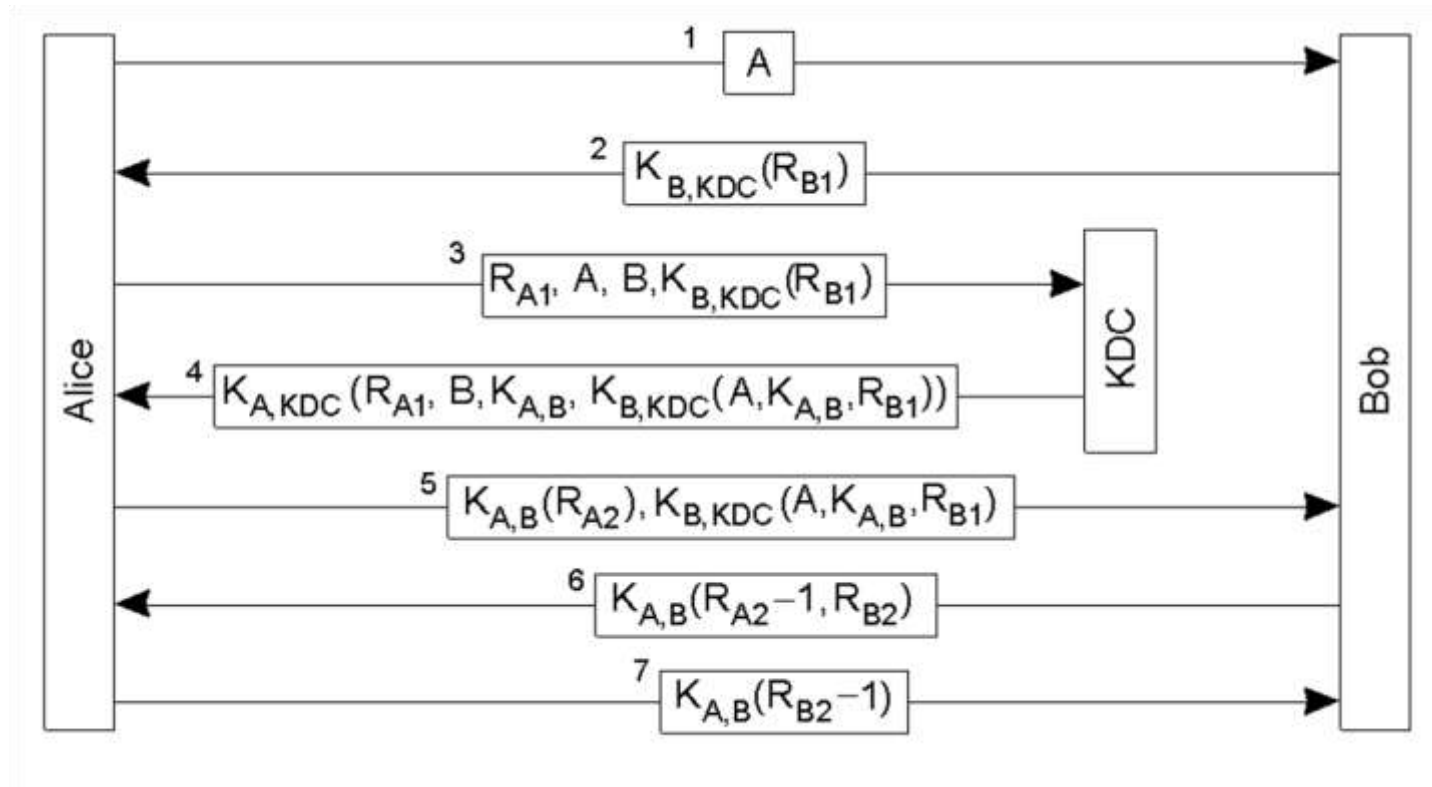
# Secure channels



- Needham-Schroeder protocol solves these issues
- A last possible attack is replay msg 3 when  $K_{A,B}$  is compromised

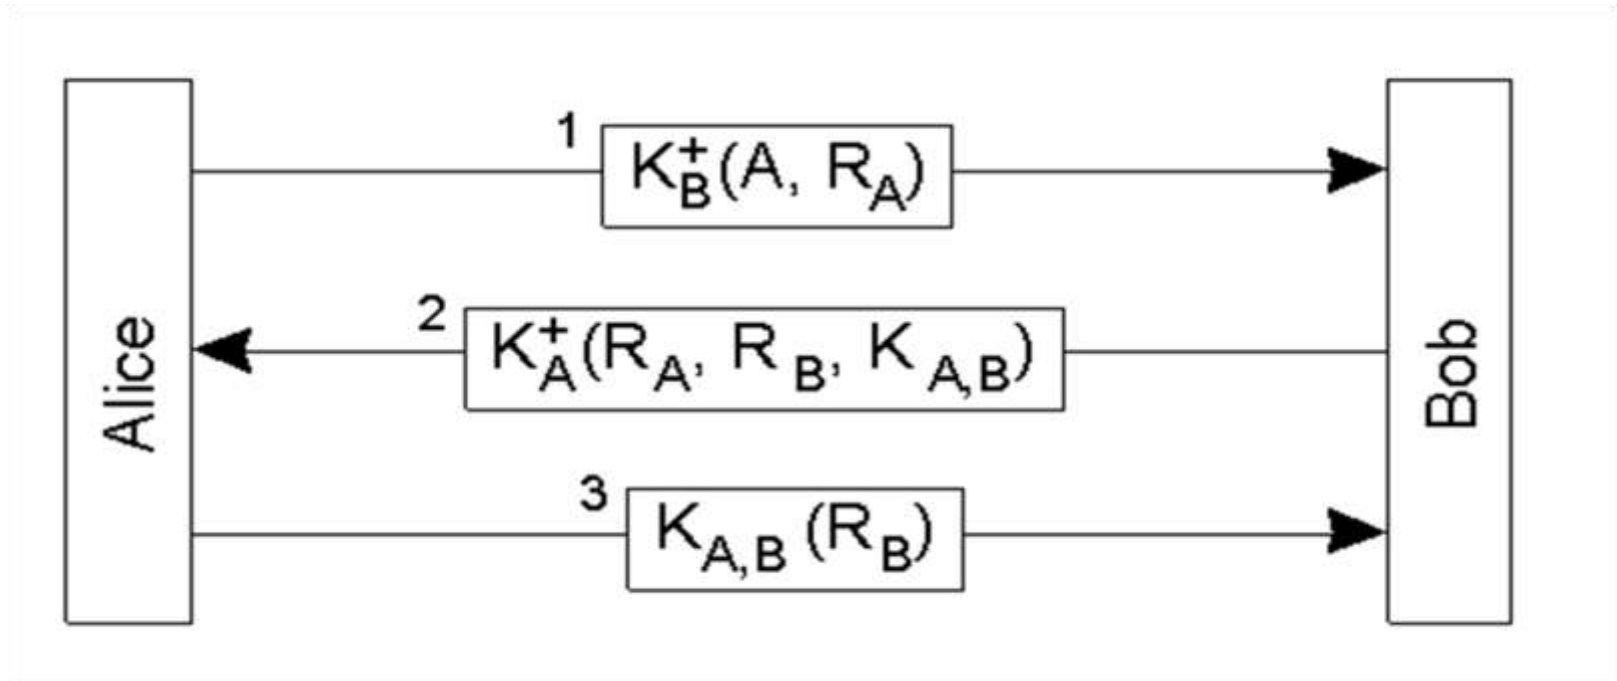


# Secure channels



- Protection against malicious reuse of a previously generated session key in the Needham-Schroeder protocol

# Secure channels



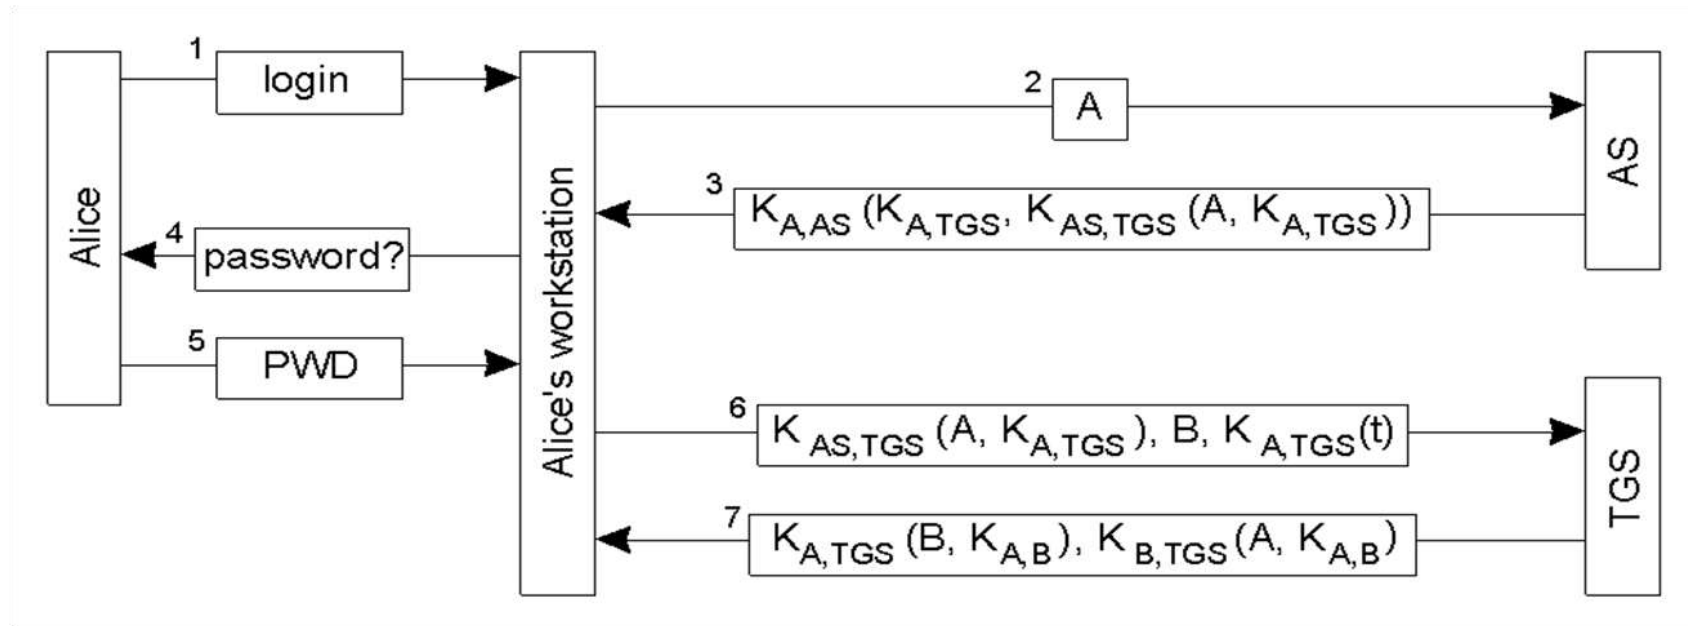
Authentication Using Public-Key Cryptography

# Kerberos

- Authentication and security facilities for intranets
  - Born in the 1980s @ MIT
  - Version 5 is a RFC
- Three security objects:
  - Ticket: to be presented to a server, granting user authentication to the system
  - Authentication: token proving user identity
  - Session key: randomly produced to be used by a client when communicating with a server (used only if needed)

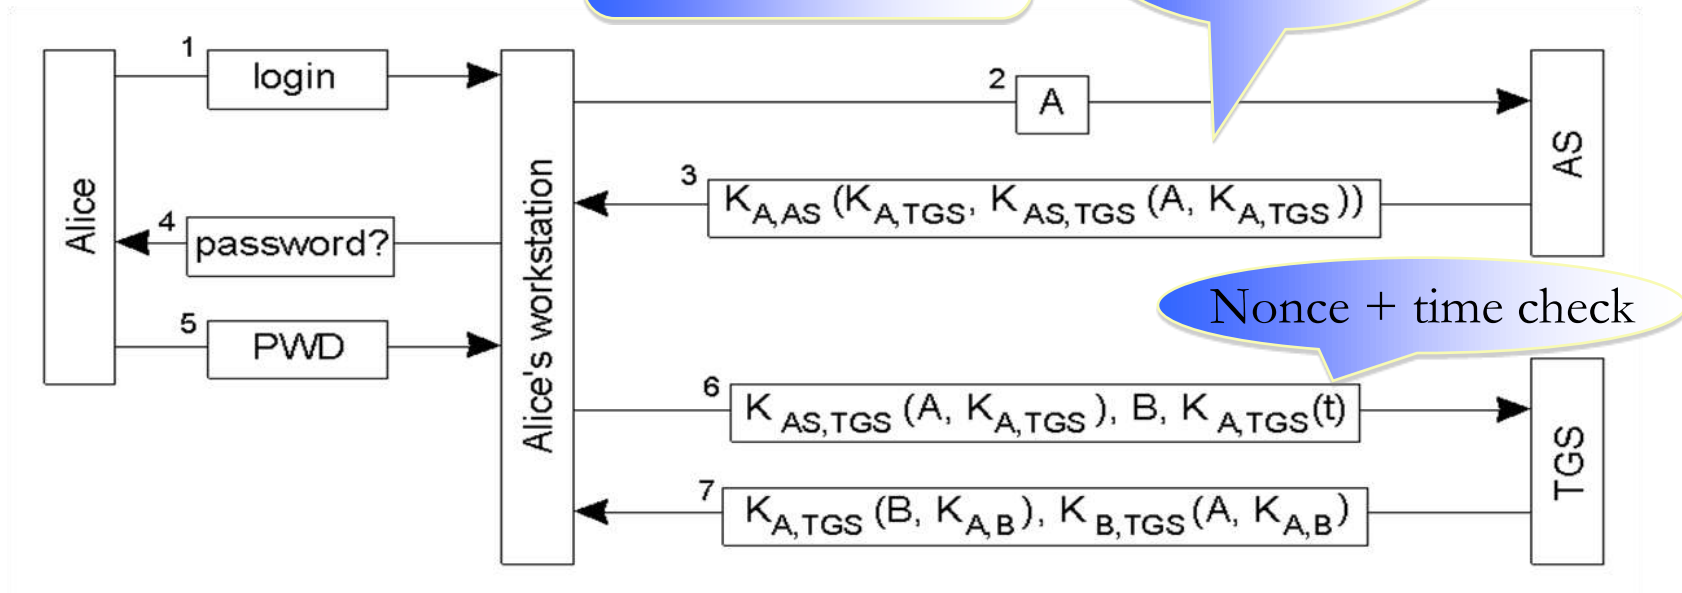


# Kerberos



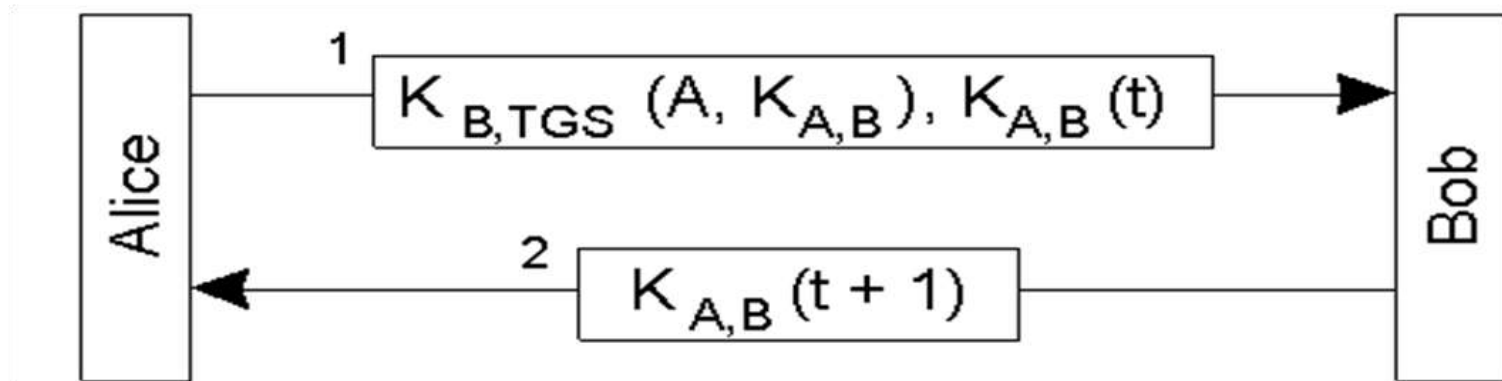
- Kerberos authentication uses Needham-Schroeder. KDC is split in two:
  - An Authentication Server (AS)
    - Communication with AS is simplified by the fact that the destination is always the TGS
  - A Ticket Granting Service (TGS)

# Kerberos



- The password is used to generate  $K_{A,AS}$  but there's no need to send it to other hosts or even to temporary store it in workstation
- Needham-Schroeder simplified because destination is always TGS

# Kerberos



- The scheme establishes what is known as single sign-on (within the time limit imposed by the ticket)
  - As long as Alice does not change workstation, there is no need for her to authenticate herself again to any other host of the distributed system
- Setting up a secure channel is easy:
  - Bob can be sure message 1 is coming from Alice because of  $K_{B,TGS}$
  - By responding with message 2 Bob proves his own identity to Alice



# Secure group communication


# Secure Group Communication

- Suppose we want secure group communication
  - Symmetric encryption with a key for each pair of participants
    - Encryption  $O(n)$ . Requires  $O(n^2)$  keys
  - Public key encryption
    - Encryption  $O(n)$ . Requires  $O(n)$  keys
    - Computationally very expensive
  - Symmetric encryption with a single shared key
    - Encryption  $O(1)$ . Requires only one key

Cost	Symmetric, pair	Public	Symmetric, single
Processing	$n$	$n$	1
Keys in the system	$(n^2-1)/2$	$2n$	1

# Secure Group Communication

- Joining and leaving a group
  - **requirement:** *Backward* and *Forward* secrecy
  - *Backward*: cannot decrypt messages before join
  - *Forward*: cannot decrypt messages after leave



- Solution: change the group key
  - Encryption of the new key
    - Join: with the old group key, with the joiner's **key:  $O(1)$**
    - Leave: with every remaining member's **key:  $O(n)$**
  - Leave is a costly operation!


These are personal keys  
not used for group  
communication

# Secure Group Communication

- Who will choose the new key?
- The server
  - Key Distribution Problem
    - We will see two examples of efficient techniques for performing key distribution
- The participants
  - Key Agreement Problem
    - We will see Diffie-Hellman for 2 parties
    - It can be generalized to more than two participants

# Efficient Key Distribution

- Logical Key (Tree) Hierarchy
  - leaves: members with keys
  - root: (Data Encryption) Key
  - each member knows (Key Encryption) Keys up to root
- Leave:
  - change all keys known by the leaving member
    - Data Encryption Key
    - change keys on **path** to root
  - diffuse efficiently the new keys
    - encrypt each key with child's
    - exploit **stable** subtrees for key distribution




$$K_5(K_{4-5}'), K_{4-5}'(K_{4-7}'), K_{6-7}(K_{4-7}'), K_{0-3}(K_{0-7}'), K_{4-7}'(K_{0-7}')$$



# Efficient Key Distribution

- Centralized Flat Table

- 2 KEKs for each bit of member ID + 1 DEK
  - ( $= 2 \cdot \log(n) + 1$  vs  $2n - 1$  in LKH)
- Each node has a key for each bit in its ID
  - Node 9=(1001) has keys  $K_{0,1}$   $K_{1,0}$   $K_{2,0}$   $K_{3,1}$  + DEK
  - If 9 leaves all **these keys** + **DEK** must be changed




	ID bit 0	ID bit 1	ID bit 2	ID bit 3
Bit = 0	$K_{0,0}$	$K_{1,0}$	$K_{2,0}$	$K_{3,0}$
Bit = 1	$K_{0,1}$	$K_{1,1}$	$K_{2,1}$	$K_{3,1}$

# Efficient Key Distribution

- Centralized Flat Table

- Encrypt DEK' with **every other** key: each member except 9 can decrypt the new DEK'
- Encrypt new KEKs with the old one and with DEK'
  - Send  $K_{DEK'}(K_{i,j}(K_{i,j}))$
  - Each member receives (and can read) only KEKs for its bits
- Collusion attack
  - Difficult to remove many participants



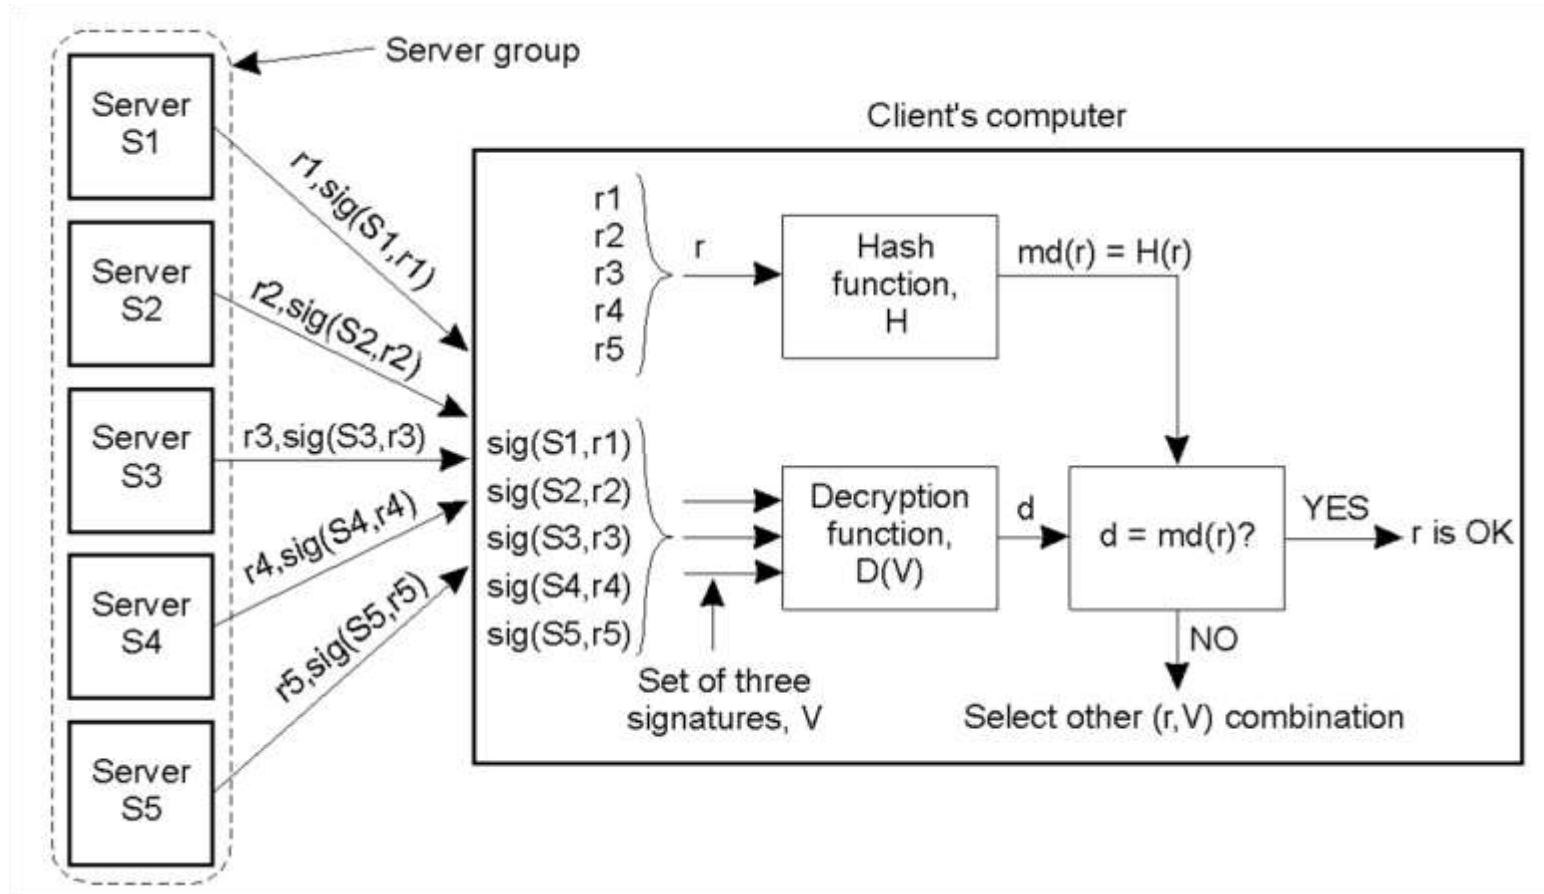
	ID bit 0	ID bit 1	ID bit 2	ID bit 3
Bit = 0	$K_{0,0}$	$K_{1,0}$	$K_{2,0}$	$K_{3,0}$
Bit = 1	$K_{0,1}$	$K_{1,1}$	$K_{2,1}$	$K_{3,1}$



# Secure replicated servers

- A client issues a request to a (transparently) replicated server. We want to be able to filter  $c$  responses corrupted by an intruder
  - First simple solution
    - Use  $2c+1$  replicated servers: each server signs its own response. The client verifies the signature and decide on a majority. However this forces the client to know the identity (and the public key) of all the servers.
- $(n,m)$  Threshold schemes
  - Divide a secret into  $m$  pieces.  $n$  pieces are sufficient to reconstruct the secret (e.g.  $n$  degree polynomial and  $m$  evaluation)
- Applied to signatures
  - Find a way such that  $c+1$  correct server signatures are needed to build a valid signature of the response
  - $r_i$  response from server  $i$
  - $\text{sig}(S_i, r_i)$  signature from server  $i$  of  $r_i$

# Secure replicated servers



Sharing a secret signature in a group of replicated servers



# Security management

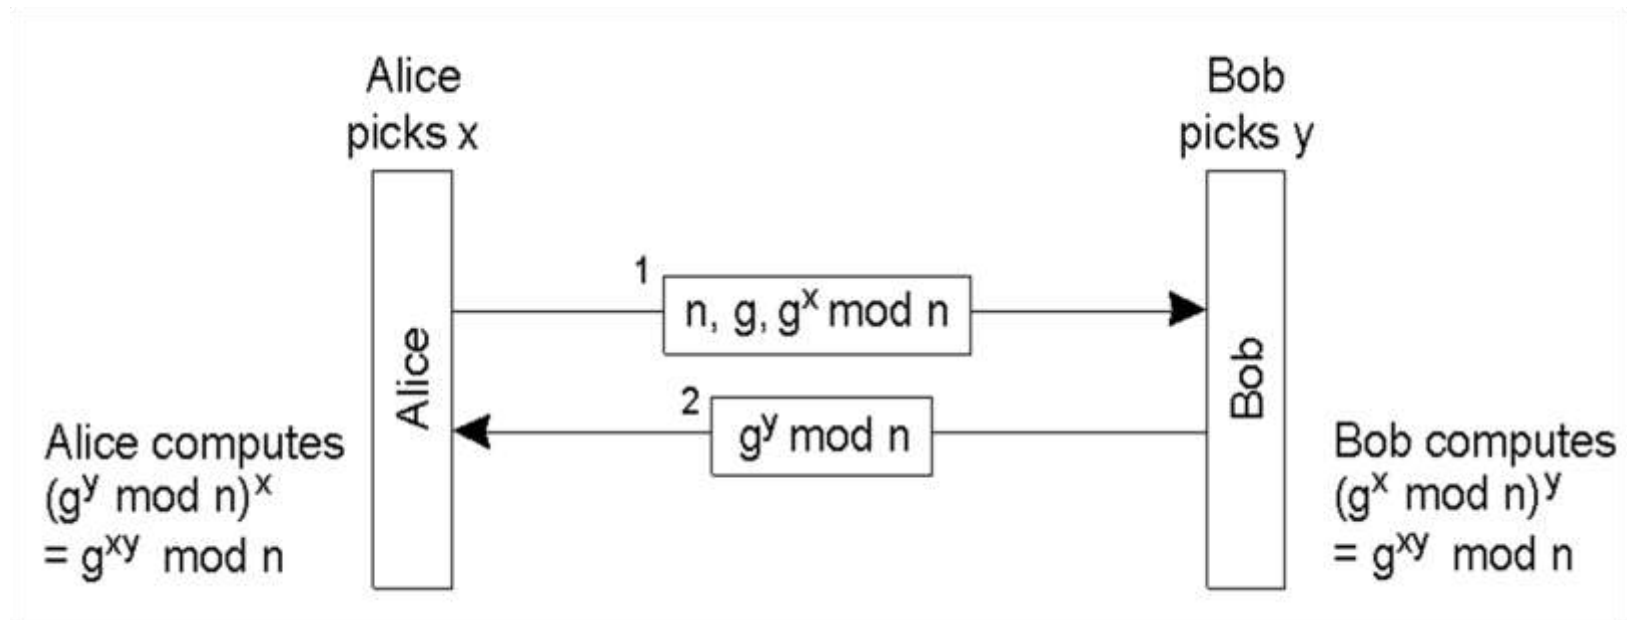
# Security Management

- We consider 3 different issues:
  1. General management of cryptographic keys
    - How public keys are distributed?
  2. Secure management of a group of servers
    - How to add or remove servers from the group?
  3. Authorization management
    - How can a process securely delegate some of its rights to others?
    - Capabilities, certificates

# Security Management

- Initial key distribution
  - We said that for authentication protocols we need one of
    - shared secret keys (between each party or with the KDC)
    - public key of the recipient of the message
  - Now we want to investigate how these initial keys can be distributed in a secure way
  - The first algorithm we consider is Diffie-Hellman key exchange
    - algorithm to exchange symmetric keys on an insecure channel

# Security Management



- Diffie-Hellman key exchange
  - $n$  and  $g$  publicly known numbers (with some good math property)
- Is this a solution?



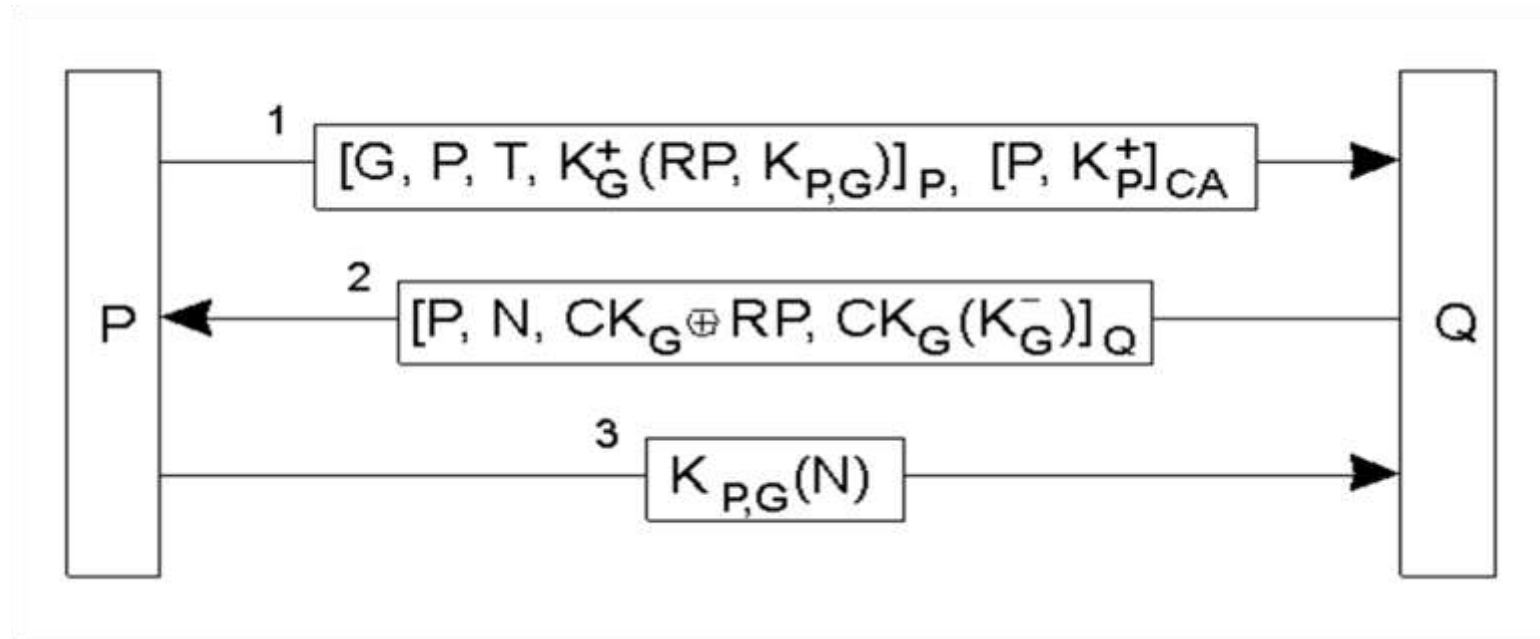
# Security Management

- Diffie-Hellman works only against passive attacks
  - If active opponent: Man in the middle attack
- We need authentication (and integrity) of both DH messages
  - These two messages can be seen as a public key exchange
- What's the usefulness of Diffie-Hellman key exchange mechanism?
  - It's a way to transform a secret key exchange in a public key exchange
  - That's good since the first requires both confidentiality and integrity, while the second only requires integrity (in particular authentication)

# Security Management

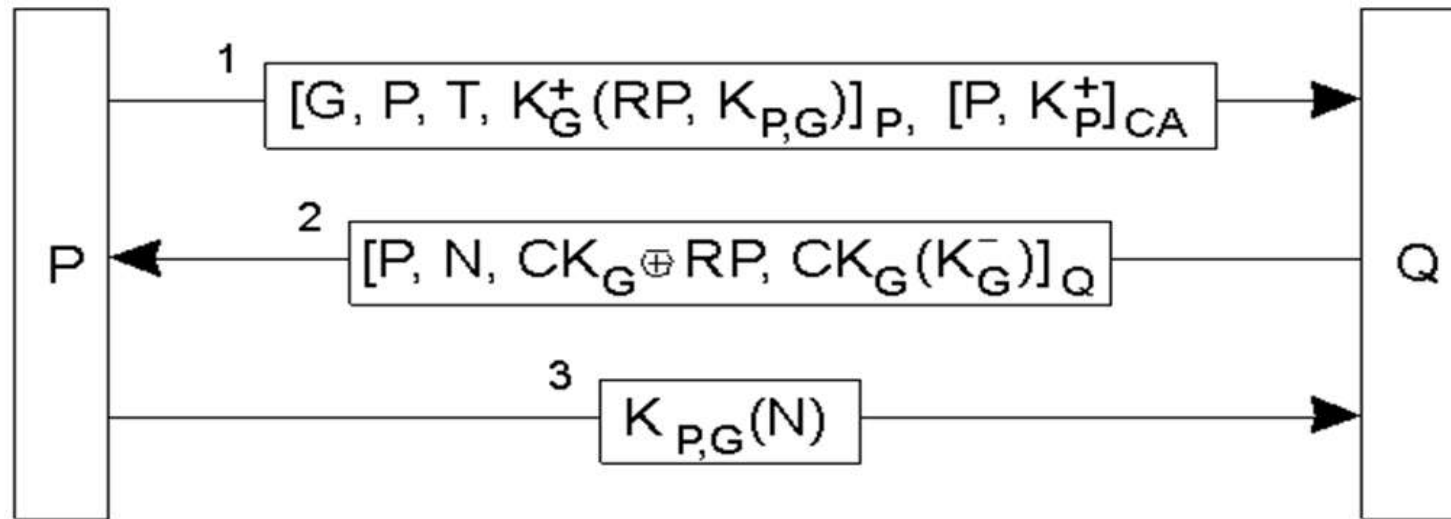
- Many systems make use of special services such as Key Distribution Center (KDC)
- These services demonstrate a difficult problem in distributed systems
  - First of all, they must be trusted
  - Second, they must offer high availability
  - Solution to high availability is replication ...
  - ... but replication makes a server more vulnerable to security attacks
- We have seen how secure group communication takes place, the problem that needs to be solved is how to manage a group of replicated servers
  - In particular, if a process joins or leaves the group, the integrity of the group should not be compromised

# Security Management



- Secure group management
- P wants to join a group G by contacting a member Q
- RP is a *Reply Pad*, generated and used only once
- Message 2 signed by Q with  $K_{P,G}$ 
  - Only a member of G could read  $K_{P,G}$  from message 1

# Security Management



- Why using RP instead of  $K_p^+$ ?
- Because RP is used only once in the second message. Using  $K_p^+$  would be risky: if  $K_p^-$  were ever to reveal, it would become possible to also reveal  $CK_G$  which would compromise the secrecy of all group communication

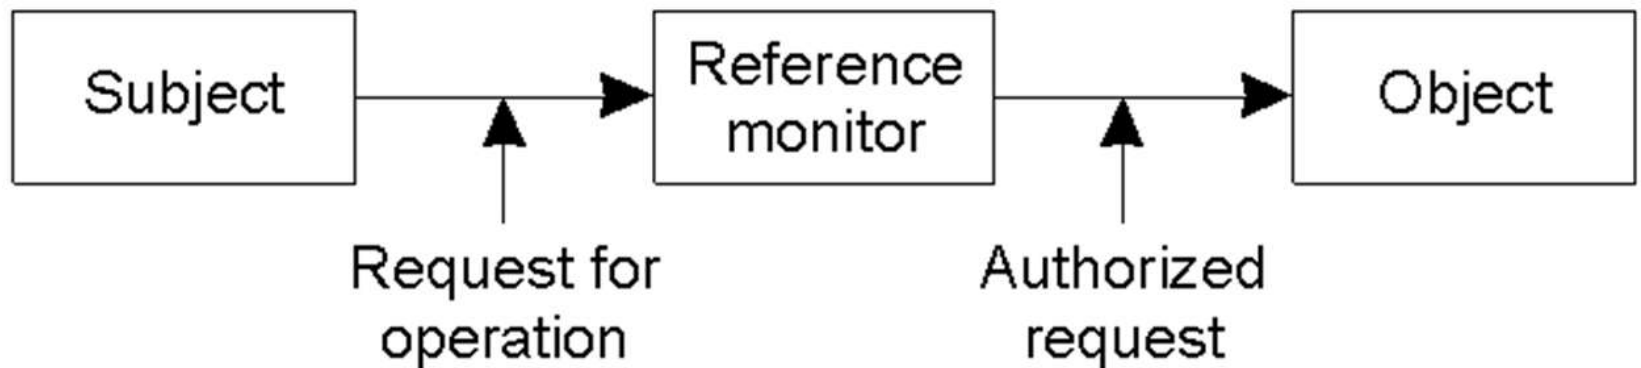
# Security Management

- Authorization Management
  - In non distributed systems, managing access rights is relatively easy
    - Each user has rights to use resources
  - In distributed systems account management is not trivial
    - We need to create an account for each user on each machine ...
    - ... or to have a central server that manages authorization rights
    - A better approach is the use of capabilities



# Access control

# Access Control



- Access control is done through a *reference monitor* which mediates requests from *subjects* to access *objects*

# Access Control

- Access Control Matrix

- Ideally we have a matrix listing rights for every combination (resource, user)

	Object 1	...	Object n
User A	R,W	...	R
...	...	...	...
User k	W	...	W

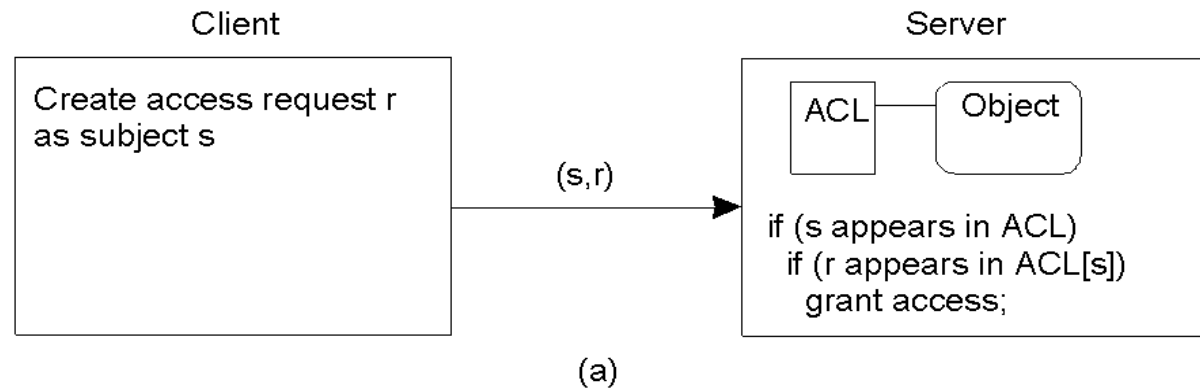
- It's a sparse matrix (we cannot implement it as a true matrix)
  - if we distribute it column-wise we have *Access Control Lists* (ACL): each object has it's own ACL
  - if we distribute it row-wise we have *capabilities lists*: each capability is an entry in the access control matrix




# Access Control

- Using ACL vs. *capabilities lists* affects the way interaction takes place with the Reference Monitor

## a) ACL

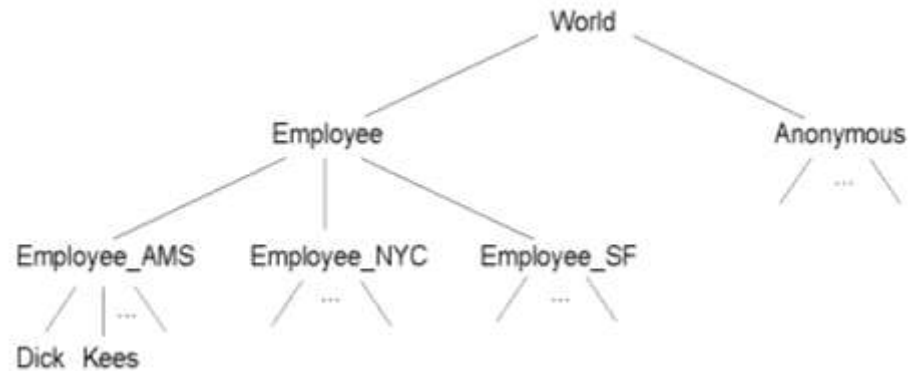


## b) capabilities



# Access Control

- ACL still occupy great memory. We can use *groups* to build a hierarchy in ACL



- Management is easy and large groups can be constructed efficiently
- Costly lookup procedure
- We can simplify the lookup by having each subject to carry a certificate listing groups he belongs to (this is similar to *capabilities*)



# Access Control

- Another possibility is role-based access control.
- Each user is associated to one ore more *roles*
  - often mapped to user's functions in an organization
- When a user logs in, he specifies one of its roles
- Roles define what can be done on which resources
  - Analogous to groups
- The difference is that users can dynamically switch from one role to another one
  - This is difficult to implement in terms of groups access control

# AC - Capabilities

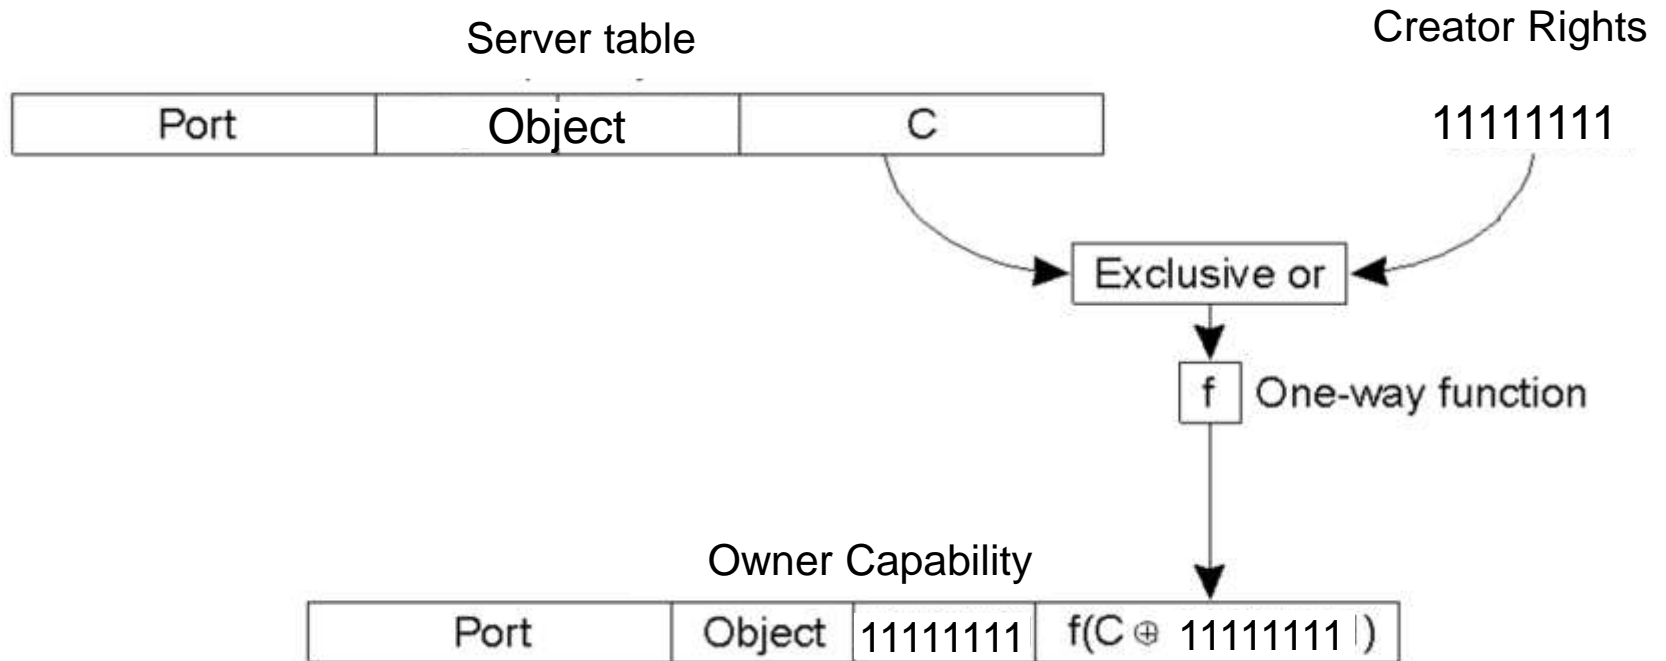
- An example of capabilities (Amoeba)
  - 128 bit

48 bits	24 bits	8 bits	48 bits
Server port	Object	Rights	Check

- 72 (48+24) to identify an object
- 8 for access rights
- 48 to make it unforgeable



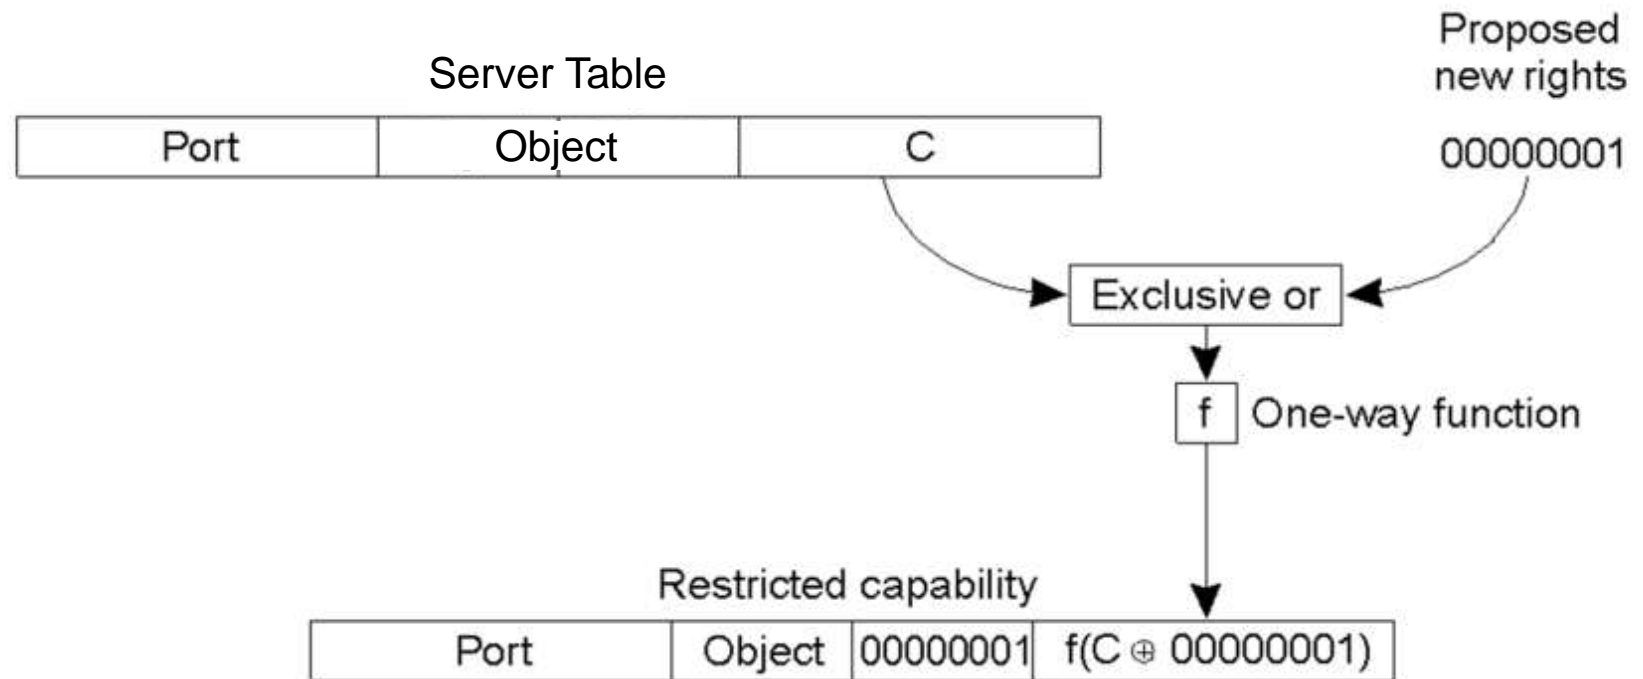
# AC - Capabilities



- On object creation, the server generates and stores **C** internally



# AC - Capabilities



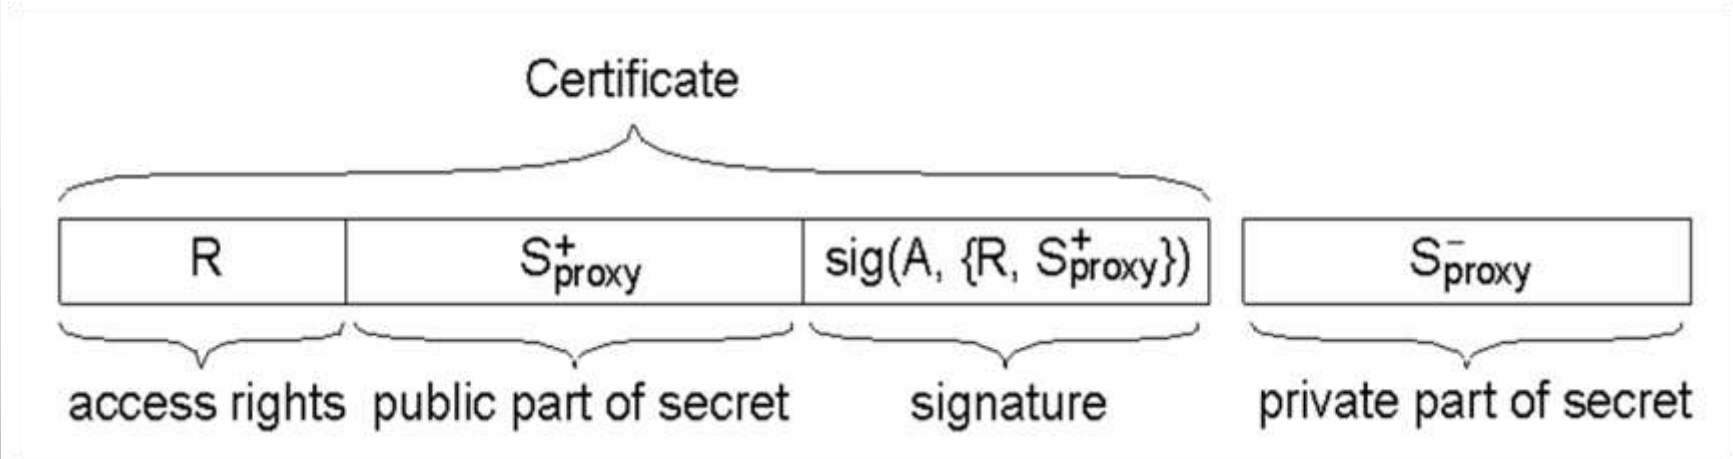
- Changing rights is not possible for the owner of the capability
  - Only the server knows  $C$

# AC - Capabilities delegation

- Delegation
  - A process may want to delegate some of its rights to other processes
    - Amoeba capabilities can only be passed as they are, it is not possible to further restrict rights unless we request a restricted capability to the server
    - A general scheme that support delegation including rights restriction is based on *proxies*
  - A proxy is a token
    - Provides rights to the owner
    - A process can create proxies with at best the same rights of the proxy they own




# AC - Proxy



- A proxy has two parts: a certificate and a key
  - The certificate proves that the *grantor* A entitled R rights to some *grantee*
  - The key is the proof that a process is the *grantee*

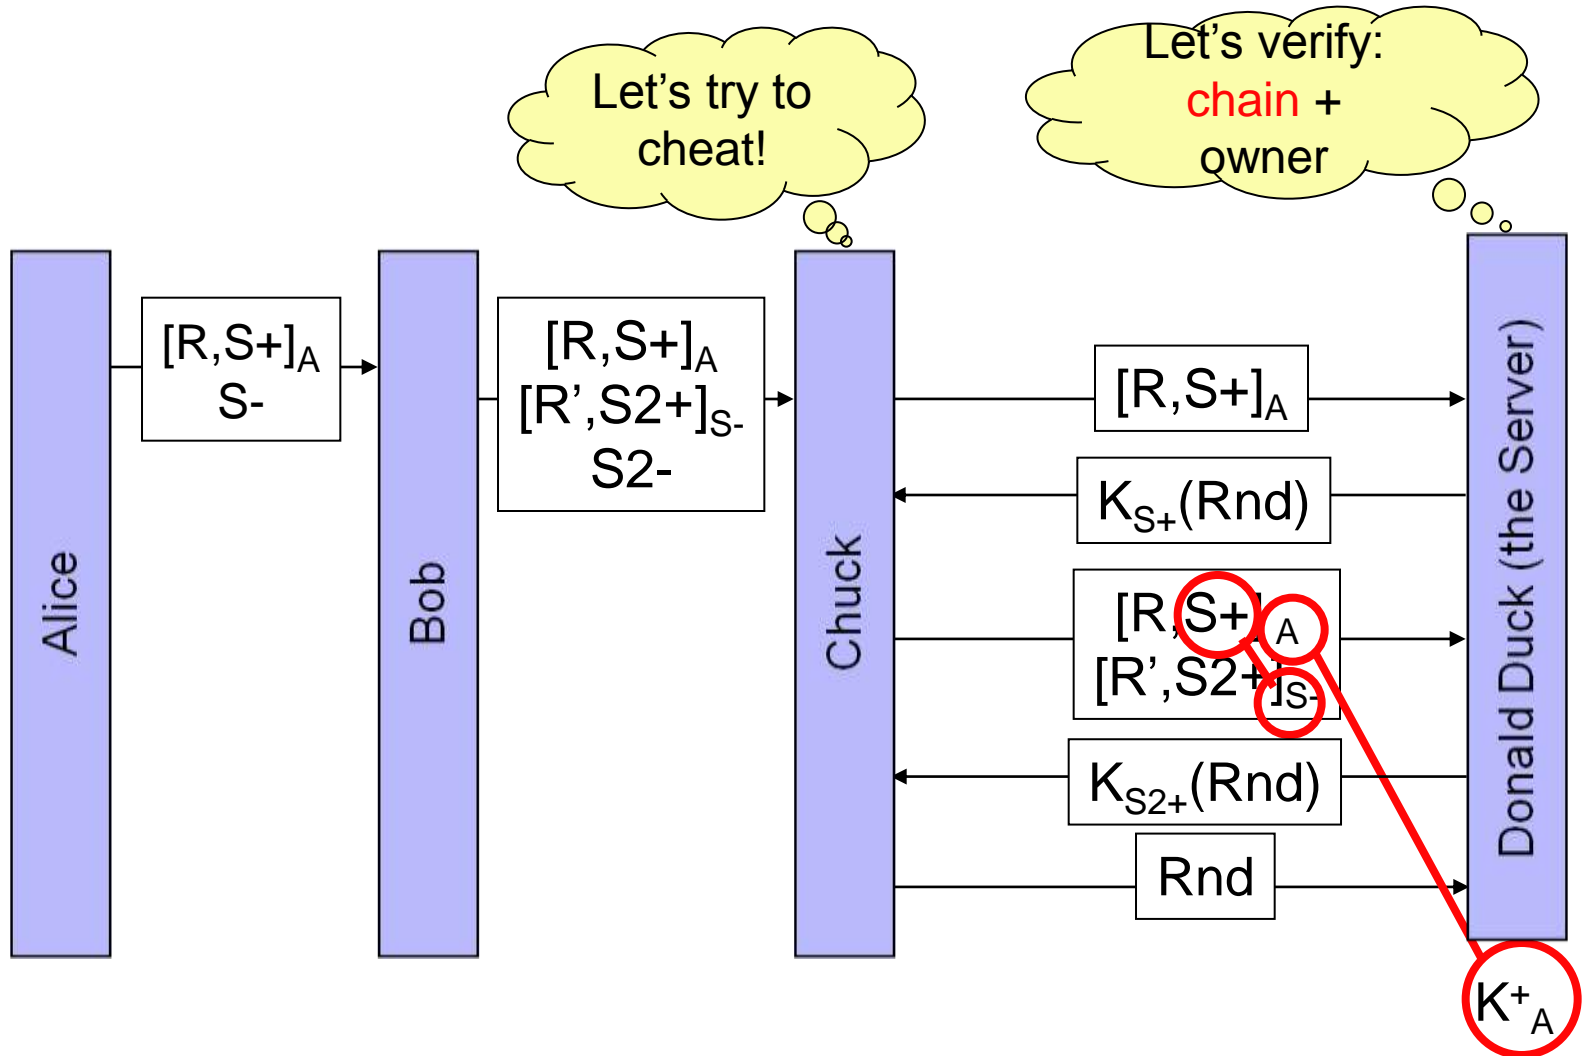


# AC - Proxy



Protocol for delegation and authorization

# AC - Restricting a Proxy



# AC - Restricting a Proxy

- How can we restrict a proxy?
- $A \rightarrow B \rightarrow C \rightarrow D$  (the server)
- B receives  $[R, S^+]_A \quad S^-$
- B places  $R_2$  restriction and transmit to C
$$[R, S^+]_A$$
$$[R_2, S_2^+]_{S^-} \quad S_2^-$$
- C can't pretend that he is entitled with R because it doesn't know  $S^-$

# Access Control - Mobile code

- Protection of the Agent
  - Agent can be altered (data and code)
  - Agent's data (possibly sensitive) can be scanned
  - Agent can be simply destroyed
- Protection of the Host
  - Typical access control issue, yet we want downloaded code to have very limited rights if compared to normal rights of the user who downloaded it

# Access Control

## Protection of the Agent

- Attack types

Leaking of code, data, control flow.	Denial of execution
Manipulation of code, data, control flow.	Leaking and manipulation of the interaction with other agents.
Incorrect execution of code.	Returning wrong results of system calls issued by the agent.
Masquerading of the host.	Tampering of agent itineraries

- Complete protection for the agent is not possible
  - at least until TCPA-Palladium or some other Big-Brother technique is massively deployed (so we wrote some years ago ☺)
- What can be done is detecting modification to an agent code or data: in Ajanta we have
  - Read only state: signed by the owner with private key
  - Targeted state: encrypted with a key not shipped with the agent
  - Append only logs

# Access Control

- Protection of the Agent
  - Append only logs: data can only be appended, not removed or modified
    - Initially the log contains only the initial checksum  $C$  (stored at the creator of the agent)
      - $C = K^+(N)$  where  $N$  is a nonce
    - Addition of  $X$  at a server  $S$ : append  $X$ , and  $\text{sign}(X, S)$
    - The new checksum  $C_{\text{new}} = K^+(C_{\text{old}}, \text{sig}(X, S), S)$

# Access Control

- Protection of the Host
  - Sandbox: each instruction can be fully controlled
    - Compiled code
      - Identify dangerous operation and shell them with code that check conditions at runtime
    - Interpreted code
      - Easier: the interpreter can control the security of each instruction execution
  - Java approach to sandbox
    - Trusted Class Loader (applets can't change it)
    - Byte code verifier (check for malformed code)
    - Security manager (a reference monitor)

# Access Control

- Protection of the Host
  - Playground (for increased security)
    - Downloaded code is diverted to a dedicated machine
    - Increased security: we have 3 level of protection
      - Sandbox
      - Playground machine OS
      - Limited connectivity from playground to other internal machines
  - Signed code
    - Security Manager can have several security policies for authenticated code



# Access Control

- Protection of the Host
  - Advanced Java techniques for security:
    - Object reference as capabilities
      - We can leverage from Java strong type system: it's not possible to forge object references
      - If the application at loading time is not given a reference to e.g. a file factory the application can't use the file system
      - Requires a redesign of the API to permit reference navigation only in safe ways
        - Java-E language
    - Name space management
      - Import statement can be resolved, based on different policy settings to different packages: `import java.io.File` could be diverted for untrusted apps to `security.File`

# Access Control - Java

- The access control architecture in Java platform protects access to sensitive resources (i.e. local files..) or sensitive application code
- All access control decisions are mediated by a security manager
  - Represented by the `java.lang.SecurityManager` class
- A security manager must be installed into the Java runtime in order to activate the access control checks
  - Java applets and web application are automatically run with a security manager installed
  - To run a local application with a security manager
    - Use the `System.setSecurityManager()` method or ...
    - ... Invoke java with `-Djava.security.manager` argument on the commandline



# Access Control - Java

- Before Java 2 (JDK 1.2) the only way to specify access control policies was to write an ad-hoc security manager object
- Local applications were always considered as trusted ones
- Since Java 2 access control architecture has changed:
  - Code from different sources can be given different sets of permissions
    - Fine-grained levels of trust
  - Access control policies can be configured in different ways
    - Es. By editing a text file
  - Also locally installed code can be treated as untrusted or partially untrusted
    - Only system classes and standard extensions are considered as fully trusted
  - Policies determine Permission objects
    - Permission subclasses can be defined to govern access to new system resources

# Access Control - Java

- How sandbox works (Java 1.0 – 1.1)
- Suppose that an applet attempts to read /etc/passwd
- First it must use a FileInputStream object passing the file name to the constructor
- FileInputStream can be a dangerous class
  - Its programmers designed its constructor to use a SecurityManager class
  - The constructor calls SecurityManager.checkRead()
  - CheckRead has no return value
    - It either returns normally or throws a SecurityException

# Access Control - Java


- How sandbox works (Java 2)
- What's new in Java 2 is how the Security Manager methods work
- In our example we described the `checkRead()` method
  - It creates a `FilePermission`
    - With target `“/etc/passwd”` and action `“read”`
  - It passes the `Permission` to the static `CheckPermission()` method of the `AccessController` class
  - One of the main features of a `Permission` subclass like `FilePermission` is the definition of the `implies()` method
    - In the case of files it can be used to determine whether permission to read `“/etc/*”` implies permission to read `“/etc/passwd”`

# Access control - Java

- So, since Java 2 we can use policy file to tune rights given to a specific application
- The class loader automatically associates the following information with the code:
  - Where the code was loaded from
  - Who signed the code (if anyone)
  - Default permissions granted to code
    - Only a limited set, all the others are determined by the access control schema seen above
- Depending on where the class was downloaded from the Security manager applies a different policy



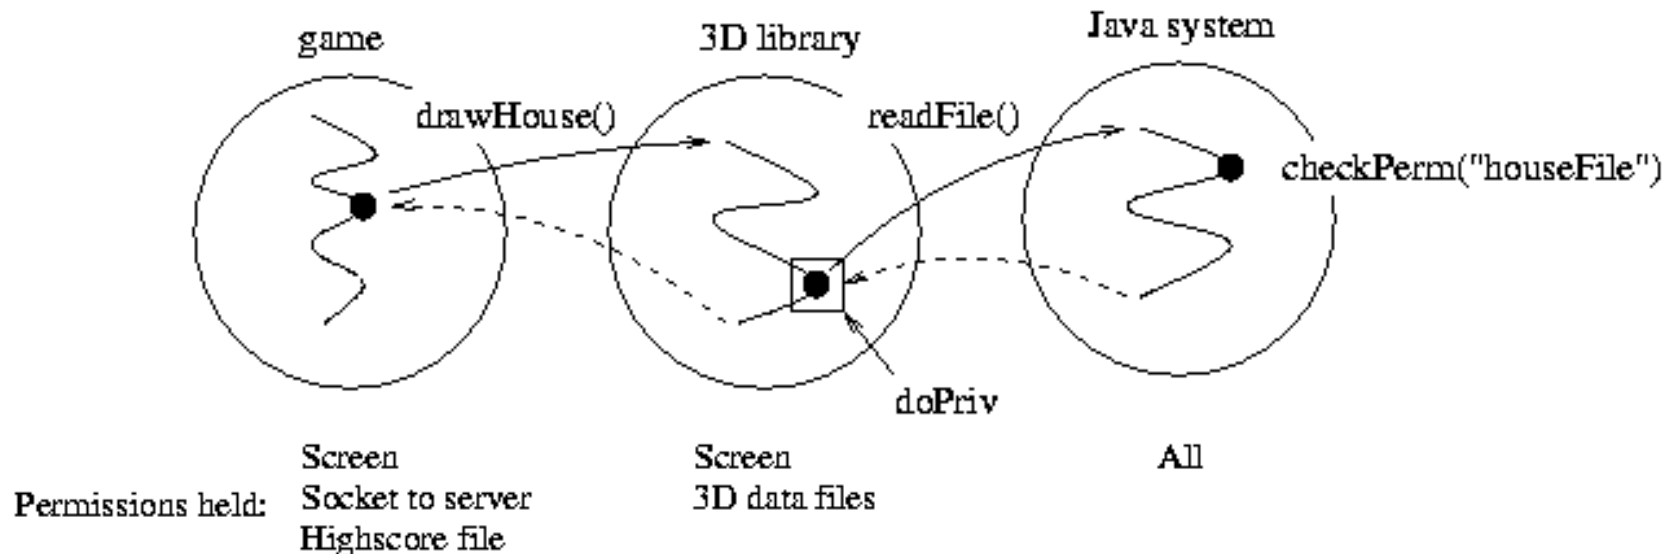
# AC - Protection of the Host



- (Extended) Stack introspection
  - Each codebase can have its own security policy
  - What happens when we have multiple procedures on the stack with different policies?
  - Default behavior: check if the call is permitted in every method of the stack

# AC - Protection of the Host

- Now 3D library wants to access 3D data files...



```
Object house = AccessController.doPrivileged(
    new PrivilegedAction() {
        public Object run() {
            return readFile("house");
        }
    });
```

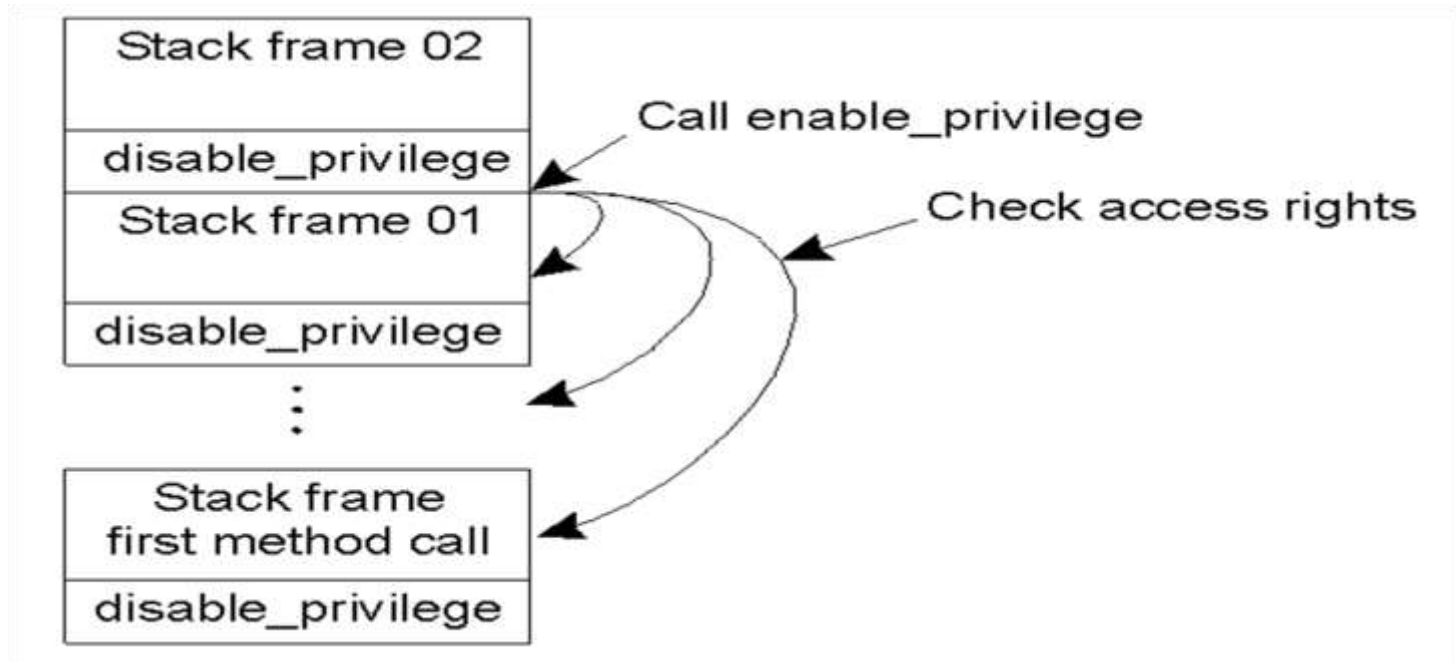


# AC - Protection of the Host

- (Extended) Stack introspection algorithm

```
Boolean checkPermission(Permission
    toCheck) {
    Stack callStack = getCallStack();
    while( ! callStack.empty() ) {
        StackFrame here = callStack.pop();
        Domain thisDomain = here.getDomain();
        if(!thisDomain.implies(toCheck)
        )return false;
        if( here.isPrivileged() ) return true;
    }
    return true;
}
```

# AC - Protection of the Host



- The old API (pre Java 1.2) used two constructs to enter and leave a privileged section
  - the book talks about these old APIs



# AC - Firewalls

- So far we have seen approaches that work fine as long as all communicating parties play according to the same set of rules.
- This happens when the distributed system is isolated from the rest of the world.
- To protect resources in an open world, a different approach is needed:
  - Any external access is controlled by a special kind of reference monitor known as Firewall

# AC - Firewalls

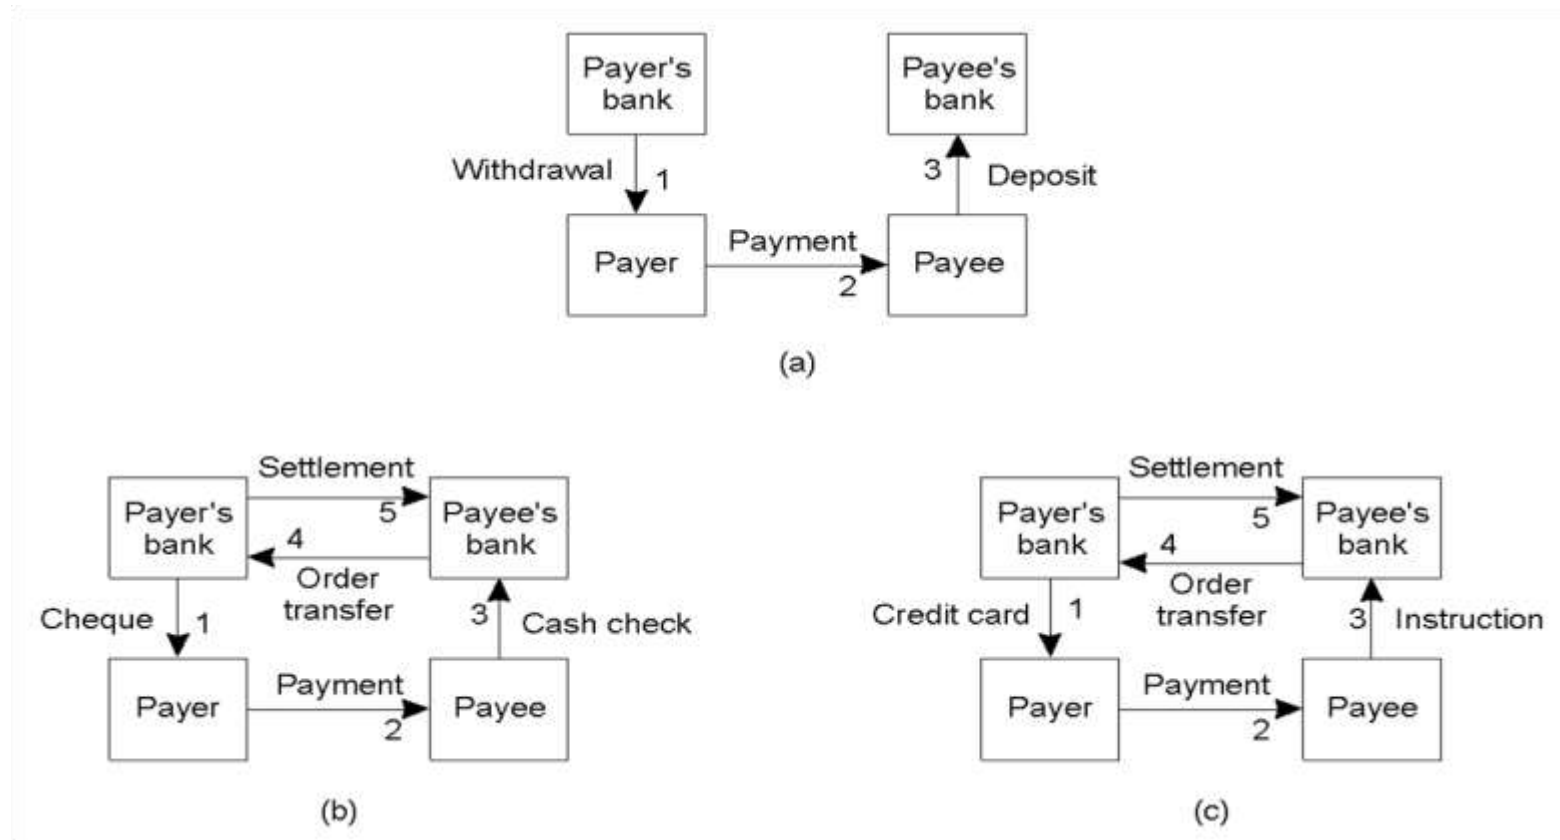
- Firewalls essentially come in two different flavours:
  - Packet filtering gateway: this type of firewall operates as a router and makes decision as to whether or not to pass a network packet based only on packet header's info (source and destination address / port)
  - Application level gateway: this type of firewall actually inspects the content of a message (example: discard incoming mail exceeding a certain size)



# Electronic payment




# Electronic Payment




a) cash b) check c) credit card

- Traditional payment methods that require co-location of customer and merchant

# Electronic Payment



(a)



(b)

a) money order b) debit order

- Other methods are used otherwise

# Electronic Payment

- There is a privacy issue

<b><u>CASH</u></b>	<b>Merchant</b>	<b>Customer</b>	<b>Date</b>	<b>Amount</b>	<b>Item</b>
<b>Merchant</b>	Full	Partial	Full	Full	Full
<b>Customer</b>	Full	Full	Full	Full	Full
<b>Bank</b>	None	None	None	None	None
<b>Observer</b>	Full	Partial	Full	Full	Full

<b><u>CREDIT CARD</u></b>	<b>Merchant</b>	<b>Customer</b>	<b>Date</b>	<b>Amount</b>	<b>Item</b>
<b>Merchant</b>	Full	Full	Full	Full	Full
<b>Customer</b>	Full	Full	Full	Full	Full
<b>Bank</b>	Full	Full	Full	Full	None
<b>Observer</b>	Full	Partial	Full	Full	Full



# Electronic Payment

- Two examples
  - E-cash
    - Provide a service similar to real cash (anonymity)
  - SET
    - Secure credit card based service
- E-cash basics
  - Alice “withdraws” cash notes signed by the bank
  - Bob can verify authenticity with bank signature
  - Bob contact the bank to protect from double spending
  - No privacy: the banks can match serial numbers

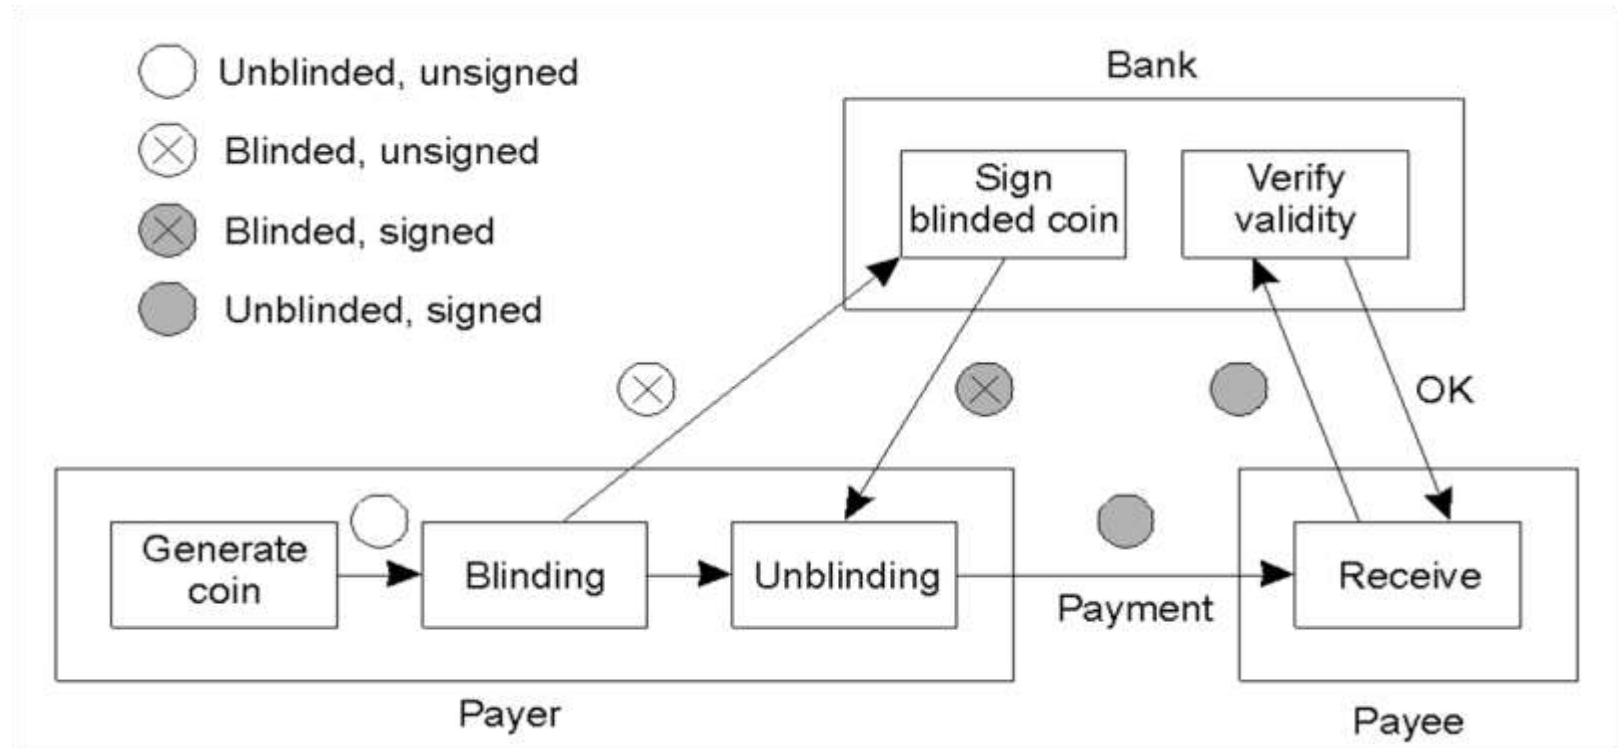
# Electronic Payment

- E-Cash
  - Blind signature
    - Alice put something in an envelope whose inside is covered with carbon paper, the bank signs it
    - The bank doesn't know what's inside but can't deny having signed it
      - Since the bank doesn't know what's inside the envelope, the amount can't be on the note. It should be in the signature. That's why the bank uses a different keys for different values (K-10 for 10€, K-1 for 1€ etc.)
    - What's inside the envelope?
      - “this is my serial number: .....”

# Electronic Payment

- Blind Signature
  - With RSA
    - e bank's public key
    - d bank's private key
    - Alice builds
      - $t = mk^e \bmod n$
    - The bank signs it
      - $t^d$
    - Alice unblinds it
      - $s = t^d/k \bmod n = m^d k^{ed}/k \bmod n = m^d \bmod n$
    - That's it: the bank signed  $m$  without knowing  $m$

# Electronic Payment



E-Cash

# Electronic Payment

- E-Cash
  - There are a number of properties that should be useful in a digital cash system
    1. Independence (from the physical world)
    2. Security (cannot be copied and reused)
    3. Privacy
    4. Off-line payment
    5. Transferability
    6. Divisibility
  - Our protocol does not satisfy 4 and 6. There are other protocols that do satisfy all of them

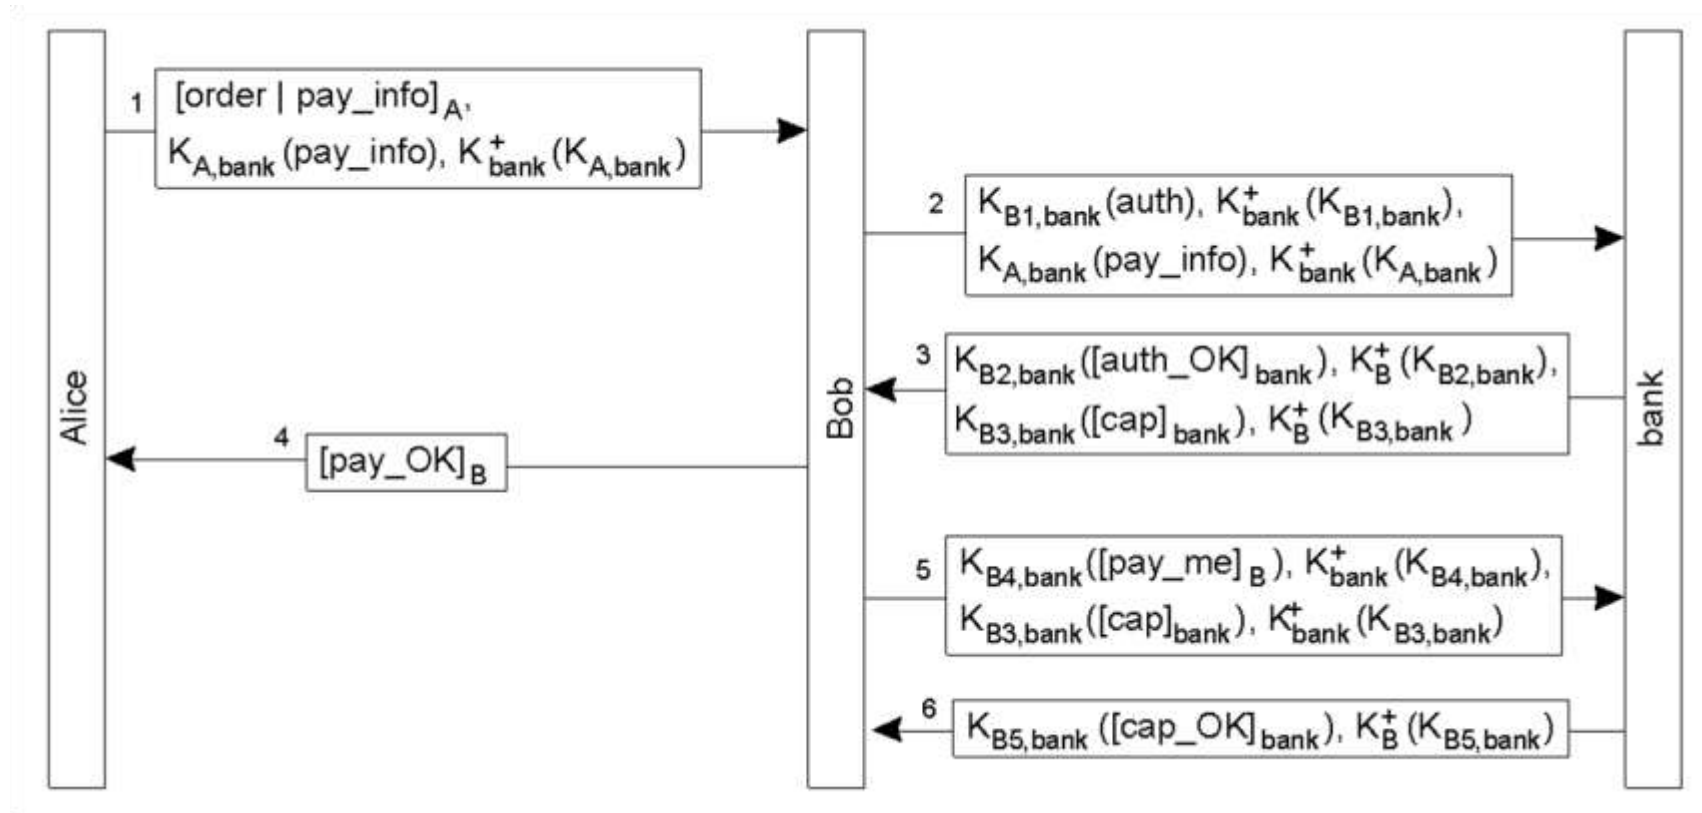
# Electronic Payment

- SET
  - We saw that credit card payments leave many e-trails.
  - Using credit card on the internet requires that the merchant does not know payment details (e.g. credit card number) and the bank does not know order details.
  - How to support credit card payments?

# Electronic Payment

- SET
  - Dual signature
  - Alice wants to send a message for approval to both Bank and Merchant, disclosing
    - the first part (m1) to Bob only (the merchant)
    - the second part (m2) to Charlie only (the bank)
  - $[m1 \mid m2]_A = [m1, H(m2), [H(m1), H(m2)]_A]$
  - $[m2 \mid m1]_A = [m2, H(m1), [H(m1), H(m2)]_A]$
  - Each approver verifies his own part and checks that it matches (with the hash of the other) the dual signature

# Electronic Payment



SET